

## Introduction

This SIF 3.0.1 Infrastructure *Read This First* (RTF) document contains a set of Frequently Asked Questions (FAQs) and their answers. Together they are designed to assist a technical readership ranging from those only casually interested in the SIF standard to those actively planning to adopt it, in quickly “getting their arms around” a secure, featureful REST-based infrastructure that took three volumes of documentation to completely describe.

This RTF serves as a non-normative “quick start” guide which does not replace the SIF specification. It reorganizes, summarizes, interprets and at times expands on the material contained in the other volumes without introducing new infrastructure requirements. As such, the RTF can be reissued as additional important questions get raised, without requiring the reissuing of a new release of the Infrastructure. It currently provides:

- An overview of the infrastructure architecture and a definition of common terminology
- A brief summary of the individual artifacts comprising the SIF 3.0 Infrastructure release (documents, schema and sandbox), and references to additional (supplementary) developer collateral.
- A set of “aspects” of the SIF 3.0 Infrastructure release (such as *Security* or *Scalability*), which define how each aspect is supported in the SIF standard, and includes specific references to where it is addressed in the SIF 3.0 infrastructure documentation.
- A work-in-progress example of a multi-staged migration strategy for a certain class of SIF v2.x deployments.

## Read this First

For those REST developers new to the SIF 3.0 standard and interested in minimizing the material they need to understand before getting started developing actual SIF 3.0 compliant software (i.e. those not willing to read this entire *Read This First* first), the following sequence of steps is suggested.

- Read chapters 1 through 3 below (the fundamentals, the architecture and the overview of existing SIF 3.0 Infrastructure artifacts and the changes introduced in SIF 3.0.1). From that point forward, everything else should make a great deal more sense.
- Go to the SIF-RS Sandbox website and get familiar with the *on-the-wire* RESTful message exchanges that occur between Consumer and Environment Provider. This “hands on” learning will make it easier to connect the later concepts to their actual implementation. The “other resources” information in section 3.5 should also be very useful to a developer trying to make a basic SIF component operational.
- Read Appendices A through C of the Utility Services document to ascertain the minimum level of functionality required of a SIF-compliant software component. This information will define the sections in the Infrastructure Services and Utility Service documents that are “must reading” for developers of any SIF 3.0 components.

Those developers already familiar with the SIF 3.0 standard should start with section 1 of the Base Architecture document, which explains the purpose of the release and references the set of changes that distinguish the SIF 3.0.1 “fix” infrastructure release from its SIF 3.0 predecessor.

# Table of Contents

<b>Introduction .....</b>	<b>1</b>
<b>1. The fundamentals .....</b>	<b>4</b>
1.1. What exactly IS the <i>SIF 3.0 Infrastructure</i> ? .....	4
1.2. Why was a major release of the infrastructure needed at this time? .....	4
<b>2. Architecture .....</b>	<b>6</b>
2.1. What are the set of special terms used to describe the infrastructure? .....	6
2.2. Are Consumers connected directly with Object Service Providers? .....	7
2.3. How does the Consumer “identify” the Service Provider? .....	9
2.4. How does the Provider “alert” the Consumer that an Event occurred? .....	10
2.5. How are Provider Responses to Consumer Requests delivered? .....	11
2.6. How can a Consumer “query” the data maintained by a Provider? .....	11
2.7. How is HTTP Utilized? .....	12
2.7.1. Message Types.....	12
2.7.2. Error Codes .....	12
2.7.3. Message Payloads.....	12
2.7.4. Non-Payload Information.....	13
<b>3. SIF 3.0 Infrastructure Documents and Artifacts .....</b>	<b>14</b>
3.1. How is the SIF 3.0 Infrastructure specification documented? .....	14
3.2. What purpose do the Infrastructure Object XML Schemas serve? .....	15
3.3. What purpose does SIF-RS (the SIF Rest Developer Sandbox) serve? .....	15
3.3.1. Learning Tool .....	15
3.3.2. REST Platform Mapping Document.....	16
3.3.3. SIF 3.0 Application Testing Aid .....	16
3.4. What purpose will the SIF 3.0 Test Harness serve?.....	16
3.5. What other resources are available for developers? .....	16
<b>4. Aspects .....</b>	<b>18</b>
4.1. How modular are SIF 3.0 Environments? .....	18
4.1.1. Environment Service “Functionality Levels” .....	18
4.2. How Secure are SIF 3.0 solutions? .....	19
4.3. How reliable are SIF 3.0 Solutions? .....	21
4.4. How Scalable are SIF 3.0 deployments? .....	22
<b>5. Migration .....</b>	<b>24</b>
5.1. Are there a set of guidelines available to end user seeking to migrate to the new release? .....	24
5.2. Is there an example of a SIF 3.0 Infrastructure Migration Strategy? .....	24

<b>Appendix A: Miscellaneous Questions .....</b>	<b>29</b>
A1. Why was there such a delay between major SIF Releases? .....	29
A2. Isn't SIF 3.0 an overly complex infrastructure specification? .....	29
A3. Why was a SIF 3.0.1 release needed?.....	30

# 1. The fundamentals

The questions and answers in this section address the overall vision of the SIF 3.0 Infrastructure release. They should be the starting point for most readers of the associated documentation.

## 1.1. What exactly IS the *SIF 3.0 Infrastructure*?

SIF 3.0 is the latest release of an open standard infrastructure which began 15 years ago as the product neutral interface of an existing commercial message broker. As will be seen, this release brings SIF into the modern era by leveraging a REST based approach to data exchange. The key contribution the SIF 3.0 release makes is to define, coordinate and standardize the ways in which a RESTful educational service can be accessed securely, robustly, and in real time by multiple RESTful clients.

SIF 3.0 infrastructure is the first infrastructure release to be completely separate from the data model defining the payloads it carries, which means it can be used to support many different data models in many different locales. For example while the SIF US data model is based on CEDS, this is not explicitly reflected in the SIF 3.0 Infrastructure documentation.

## 1.2. Why was a major release of the infrastructure needed at this time?

The SIF 2.x infrastructure was initially architected more than a decade ago, and the SIF 3.0 infrastructure introduces a wide range of needed new functionality. Of particular note are three ground-breaking design advances which satisfy long standing requests from SIF 2.x developers and implementers. They are summarized below.

### **Make the SIF Infrastructure independent of the SIF Data Model**

All current data model dependencies of the earlier release have been removed. As a result, the SIF 3.0 infrastructure can carry SIF object data from any locale (US, UK, AU), or for that matter data in conformance with other major data standards, without change.<sup>1 2</sup>

### **Offer alternatives and extensions to the required monolithic middleware component (ZIS)**

This issue has been addressed on two fronts.

First, the ZIS-provided message broker functionality has been broken up into a set of multiple, separately implementable Infrastructure Services. These were designed to individually map to common industry technologies such as “Enterprise Service Bus (ESB)”, “Queue Manager” and “Service Registry”. The SIF 3.0 *Brokered Environment*<sup>3</sup> middleware can still be implemented by a single, monolithic component, but it no longer has to be.

Second, the new SIF 3.0 architecture makes it possible to construct and deploy SIF-compliant solutions in a *Direct Environment* without utilizing any middleware at all! This new SIF infrastructure alternative offers clients a standardized

---

<sup>1</sup> Yes it's true. Footnotes in a Read This First! However they will be limited to referencing other parts of the documentation, or to provide information of interest to only a select group of readers.

<sup>2</sup> The SIF Association is currently working with the other major data standards to enable alternative data models to work with SIF 3.0 infrastructure.

<sup>3</sup> For the definitions of both Brokered and Direct Environments, please refer to the Glossary provided in the Architecture section of this document.

subset of the functionality available from the *Brokered Environment*, which means that these clients can be deployed into middleware-centric solutions with no recoding or reintegration efforts required.

### **Support industry standard transport technologies**

The SIF 3.0 infrastructure documentation describes the service framework and associated core services and utilities in a platform neutral manner. As a result, it can be mapped to any modern transport running over HTTP/S.

The defined platform mapping of the SIF 3.0 infrastructure is REST. The SIF 3.0 infrastructure includes paged reads, eTags, synchronous IO, pure data payloads and support for the other primary REST resource design patterns.

XQuery scripts are also supported in SIF 3.0 replacing the earlier SIF-specific (and less powerful) *Query* and *Extended Query* functionality.

These and other changes allow SIF 3.0 solutions to be deployed in the Data Center of an educational organization using the identical technologies that are already present and known to IT personnel. It also makes it easier for vendors to staff SIF-related projects as both the REST infrastructure technology and to a lesser extent, XQuery scripting are already familiar to large numbers of today's software developers.

## 2. Architecture

The most effective path to acquiring a working knowledge of the SIF 3.0 infrastructure starts with gaining an understanding of the architecture which guided its construction.<sup>4</sup>

This section provides a brief overview of that architecture including:

- Underlying terms and concepts
- Major infrastructure “components”
- Basic functionality
- Supported Message Exchange Patterns (MEPs)
- HTTP Usage

It is the longest section in this RTF.

The reader should have at least a passing understanding of the information presented below before tackling the three volumes of infrastructure documentation.

### 2.1. What are the set of special terms used to describe the infrastructure?

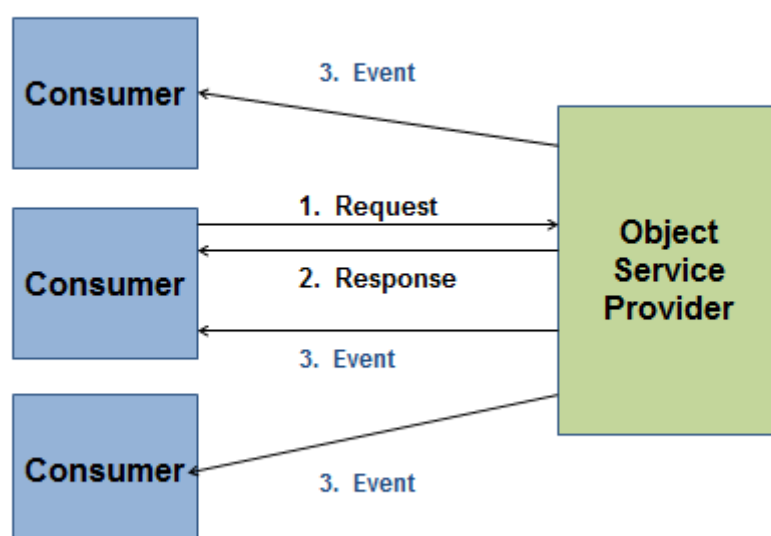
The following entries are abbreviated (rather than rigorous) descriptions selected from the “Glossary of Terms and Concepts” contained in the Base Architecture volume. They comprise the basic terminology – the terms used in the definition of other terms.

Term	Meaning
<b>Object Type</b>	Similar to a “Class” in an object oriented programming language. Every Object Type is defined by its corresponding XML Schema which determines the format of each of its internal elements.  Ex: student
<b>Object</b>	An “instantiation” of an Object Type possessing a unique ID, which can be conveyed “over the wire” and validated in accordance with its XML schema.  .Ex: student with ID 12345 (John Jones)
<b>Service Interface</b>	An interface supporting one or more of the standardized CRUD ( <i>Create</i> , <i>Query (READ)</i> , <i>Update</i> and <i>Delete</i> ) data requests on objects of a specific type.  Ex: Student Data Object Service
<b>Event</b>	A message reporting a data change to one or more objects of the same type, often published in response to a successful <i>Create</i> , <i>Update</i> or <i>Delete</i> data request.  Ex: Student 12345 has updated home phone number: 555-789-4321,
<b>Object Service</b>	An application providing access to Objects of a selected type in accordance with the Service

<sup>4</sup> All of the information in this section is provided in greater detail in the SIF 3.0 Infrastructure documentation, and particularly the SIF Base Architecture volume.

<b>Provider</b>	Interface, and publishing related Events Ex: The Student Object Service Provider
<b>Service Consumer</b>	Any application which makes requests of, subscribes to, and receives Events from, one or more Service Providers. SIF 2.x Equivalent: Client

## Consumer – Provider interactions



### 2.2. Are Consumers connected directly with Object Service Providers?

No. There is a third component in the middle, called an Environments Provider, which connects the Service Consumers to the Object Service Providers through a set of *Infrastructure Services* which taken together, comprise the Environments Provider interface. Environments are the way “client sessions” are managed, and reflect a considerably more sophisticated approach than traditional REST interfaces. It is the Environment interface that standardizes exactly how SIF 3.0 data exchanges are guaranteed significant levels of security, enhanced scalability and flexible functionality.

A Consumer’s Environment is made available when it initially registers with the Environments Provider Service. It comprises the totality of every possible service the Consumer might ever access. Depending on the authentication provided by the Consumer at registration time, the Environment returned might be testing or staging rather than actual production.

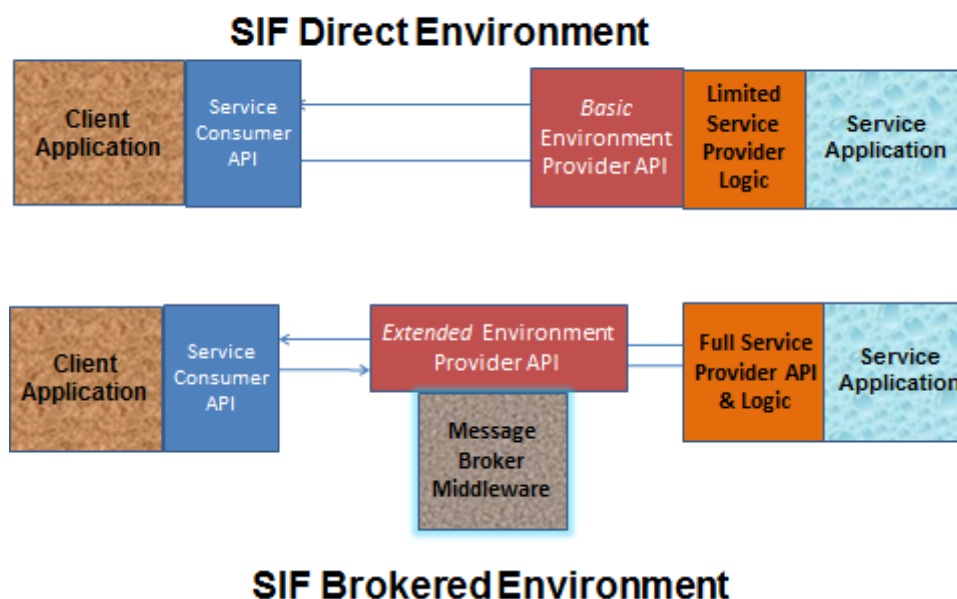
The physical implementation of the Consumer’s Environment interface can take one of two forms.

Environment Type	Implementation Topologies
<b>Direct</b>	Introduced to the SIF architecture with this release, the Direct Environment standardizes SIF-

	<p>compliant message exchanges between Consumer and Provider in the absence of a central Message Broker</p> <p>The Direct Environment connects a <u>single</u> Consumer to a fixed set of one or more directly accessible Service Providers. All Consumer to Provider connections are direct (no middleware components), because the Environments Provider Interface and all Service Provider Interfaces are implemented by a single application (such as an SIS, LMS or Data Store).</p> <p>Such an application generally provides a separate Direct Environment to each of several Service Consumers, to enable them to directly access and update the object data it provides. In that common case:</p> <ul style="list-style-type: none"> <li>• Each Service Consumer in a Direct Environment is operating independently of all other Consumers.</li> <li>• When any Service Consumer request causes a change to the object data maintained in the Service application, every appropriately subscribed Service Consumer receives the identical Event.</li> </ul> <p><b>Example:</b></p> <p>The typical Service Consumer registered in a SIS-provided Direct Environment could be a simple data entry application running on a mobile device, or a Student Contact system that only needed to access the Student's ID, Name, Addresses and Phone Numbers.</p>
<b>Brokered</b>	<p>The Brokered Environment securely and reliably connects N Service Consumers to a dynamically changing list of M Service Providers through a Message Broker.</p> <p>Unlike the Direct Environment, any Service Consumer with the proper authorization rights can provision itself as a Service Provider, and receive Requests from and publish Events to, other Service Consumers with the appropriate authorization rights.</p> <p>A Brokered Environment will commonly supply Consumers with access to services providing additional functionality such as the ability to self-provision against a list of provided services.</p> <p>The "Message Broker" functionality requirements of a SIF 3.0 Brokered Environment can be implemented (among other alternatives) by SIF "business logic" layered on top of an Enterprise Service Bus (ESB), by internally coupled middleware components or by an upgraded SIF 2.x Zone Integration Server (ZIS).</p> <p>One major feature of this architecture is that the Brokered Environment offers a superset (rather than replaces) the functionality of the Direct Environment. As a result, <i>any Consumer application which interoperates successfully in a Direct Environment can be redeployed into a Brokered Environment without reprogramming.</i></p> <p>SIF 2.x Brokered Environments Provider equivalent: ZIS</p>



## A Comparison of Environment Types



### 2.3. How does the Consumer “identify” the Service Provider?

Content-based routing is the ability to route a Request to the appropriate Service instance based upon a specified “service identification”. If the Consumer specifies that ID, the Request will be successfully delivered. In order for this to work, the identification of each Service must be unique.

In the SIF 2.x Zone, every supported object type can have only a single Object Provider. Wherever there are two systems supplying the same type of data (ex: an SIS and a Special Ed application both provide student data within a school), SIF administrators must create two separate Zones. This approach allows the name of the object type provided by a Service to completely identify that Service within the Zone, but often requires a complex N:N mapping of applications to Zones in order to adequately reflect the needs of the educational organization.

With SIF 3.0, the “Zone” Service Scoping of SIF 2.x is extended both upward (Environment) and downward (Context). Service Identity within an Environment now includes the Zone, the Object Type and the Data Context.

These additional components of this more flexible service identity are individually described in the table below.

Service Scope Limit	Meaning
<b>Environment</b>	The Environment provided to a Consumer (whether Direct or Brokered) can be divided into multiple Zones, each created by the Environment Administrator to correspond to a physical component within the owning educational organization.
<b>Zone</b>	A Zone scopes a collection of Application Services within the Consumer’s Environment. A given Service implementation may support the same Provider interface in several Zones.

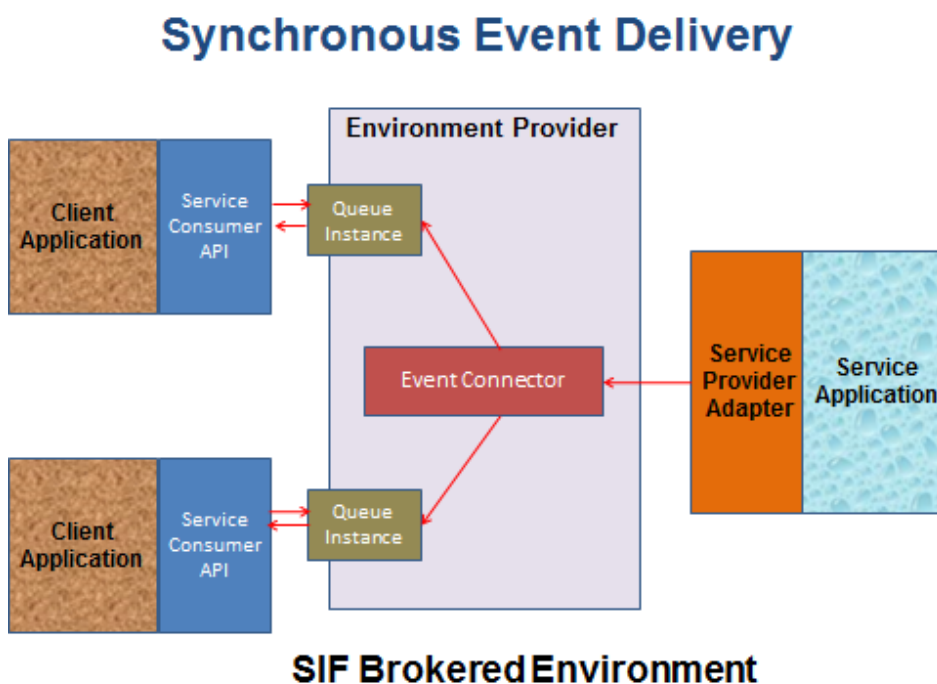
<b>Context</b>	<p>A Context is optional Data Model-specific metadata that may accompany a Consumer Request as a way of further scoping and restricting the selected Provider. For example a supplied Context might indicate that the Student Schedule Provider Service being requested in Zone XYZ is the one dealing with next term's data, rather than the current one.</p> <p>A Zone can contain multiple Object Provider Services, each offering its data in a differently named context, or a given Service implementation may support multiple contexts for the same Provider interface.</p> <p>There can be only one Service of a given Object Type with a given Context in any one Zone.</p>
----------------	---

## 2.4. How does the Provider “alert” the Consumer that an Event occurred?

The publishing of a Service Event (reflecting a change to one or more data objects) is asynchronous from the point of view of a subscribing Consumer since the Service Provider can issue it at any time. The reception of a Service Event is synchronous however. It is deposited into the message queue selected by the Consumer for that purpose when it initially subscribed to change Events for that Object Type, and the Event is later retrieved in response to a Consumer Request to the corresponding Queue Instance Service.<sup>5</sup>

The Service Provider remains unaware of the identity of the particular Consumers who might be subscribed to the Events it issues, and in fact operates independently of the number of subscribers, or even of whether there are any subscribers at all. It publishes all its Events to the Event Connector Service<sup>6</sup>, which is responsible for multiplexing the Event message and dispatching a copy of it to the Queue Instance Service of each subscribing Consumer.

The Event delivery process is illustrated in the diagram below.



<sup>5</sup> For a more complete description of the Queue Service please refer to the appropriate section in the Infrastructure Services document.

<sup>6</sup> For a more complete description of the Event Connector Service please refer to the appropriate section in the Infrastructure Services document.

## 2.5. How are Provider Responses to Consumer Requests delivered?

In addition to Events, the Queue Instance may also contain delayed Responses to previously issued Consumer Requests. The Consumer specifies in the Request whether the Response is to be immediate or delayed.

The details of each mode are described in the table segment below, which is replicated from a larger table in the Base Architecture document.

Response Mode	Details
<b>Immediate</b>	<p>The Response to a Request is provided synchronously in the immediate HTTP response, and the Requester thread for that connection “blocks” until the Response arrives.</p> <p>Immediate Responses are new to SIF 3.0. They match the standard RESTful Client design pattern and are supported in both Direct and Brokered Environments.</p>
<b>Delayed</b>	<p>The Consumer issues the Request which is replied to with an “Accept” status code in the immediate HTTP response, which indicates “<i>Request is legal and can and will be delivered to the indicated Service Provider</i>”. This frees a single-threaded Consumer to do other things.</p> <p>The Response issued by the Service Provider arrives asynchronously at a later time, in a manner identical with that of an incoming Event. It contains the Message ID of the original Request it completes.</p>

## 2.6. How can a Consumer “query” the data maintained by a Provider?<sup>7</sup>

There are several options for the developer of a SIF 3.0 Service Consumer to choose from. The first is “unqualified” Query, where the Query Responses return all elements from all objects associated with the Query URL.

Query Options	Response
<b>Bulk</b>	<p>All the data for all the objects of the specified type is returned from the Service Provider. As the amount of data may be too large to pack into a single multi-object Response message (causing a Response Error to be generated instead), the Paged alternative shown below is often selected instead.</p>
<b>Paged</b>	<p>A “Paged” Query is typically one in a series of Service Consumer queries for object data which defines the start and end of the particular Page of results from the Query.</p> <p>Example: A teacher holding a tablet device runs a simple Service Consumer REST application which interactively queries the Assessment System for the first 30 Student scores, which are returned immediately.</p> <p>After they are displayed, the teacher may hit “next” whereupon a new interactive Query will be issued, and the next 30 Student scores will be immediately returned and displayed.</p>

There are also three separate and powerful ways to issue Query Requests which impose “constraints” on the object data (Bulk or Paged) returned in the Query Response:

<sup>7</sup> Please refer to the Query section in chapter 5 of the Base Architecture Guide for a more complete explanation.

- Use Dynamic Query to specify the “where” criteria which every returned object has to meet (ex: “*Return every Student who is a senior*”). See section 5.7 of the Base Architecture document.
- Use Service Paths to selects the occurrences of one object type in another (ex: *Return every Student in School XYZ*). See section 5.5 of the Base Architecture document
- Use powerful predefined XQuery Templates to select exactly which object elements within each qualifying object of whatever object type (including calculated aggregates) are to be returned. See section 5.56 of the Base Architecture document

## 2.7. How is HTTP<sup>8</sup> Utilized?

The REST “transport” is for the most part a set of conventions about how to utilize the underlying HTTP line protocol. SIF 3.0 adopts these conventions and expands them where necessary.

### 2.7.1. Message Types

The SIF 3.0 Request / Response types map directly to the common REST usage of the HTTP line protocol:

SIF 3.0 Request / Response Type	HTTP Message Type
Query	GET
Create	POST
Update	PUT
Delete	DELETE

### 2.7.2. Error Codes

The SIF Error message codes map directly to the HTTP Error codes.<sup>9</sup>

### 2.7.3. Message Payloads

The payload of every SIF 3.0 HTTP message is XML formatted, and contains some or all of the elements of one or more Data Objects. Specifically

HTTP Message Type	SIF 3.0 Message Type	SIF 3.0 Message Payload
GET	Query Request	Empty.
	Query Response	One or more data objects matching the type of Query Request issued
POST	Create Request	One or more data objects, whose elements (excluding the ID, which is assigned by the Service Provider) are provided by the Consumer.
	Create Response	A set of “results” indicating the status of the create for each object
	Create Event	One or more data objects successfully created. The entire object as

<sup>8</sup> Actually it is more correct to ask how HTTPS is used, as that is the primary line protocol for SIF 3.0 due to security concerns. However HTTP is optionally supported, and for reasons of clarity will be the protocol referred to.

<sup>9</sup> Please refer to Appendix C of the Infrastructure Services document for the complete list of supported HTTP Codes.

		known to the Provider will be included (i.e. everything returned but non-supported optional elements)
	Update Event	One or more data objects successfully updated. Only those elements actually updated are included, unless the “replacement” flag is set in the HTTP Header, indicating that all the elements are included, whether modified or not.
	Delete Event	The IDs of one or more data objects successfully deleted.
PUT	Update Request	One or more data objects, including the ID, and containing only those elements whose values are to be modified.
	Update Response	A set of “results” indicating the status of the update for each object
	Multi-object Delete Request	Since REST conventions do not support payloads on HTTP DELETE messages, all multi-object Delete Requests are conveyed via an HTTP PUT message containing an additional HTTP Header Field value of <i>methodOverride</i> set to <i>DELETE</i> .
	Multi-object Delete Response	A set of “results” indicating the status of the delete for each object
DELETE	Single Object Delete Request	The ID of the single data object to be deleted
	Single Object Delete Response	The “results” indicating the status of the deletion for that object

#### 2.7.4. Non-Payload Information

The “non-object” data associated with these messages (ex: Page Number and Page Size of Query, or Zone and Context of the selected Service Provider) are conveyed in a variety of “non-XML” ways<sup>10</sup>.

- Segment in Service URL (Requests only)
- URL Matrix Parameter (Requests only)
- URL Query Parameter (Requests only)
- HTTP Header Field (Requests, Responses and Events)

Every *Resource* (ex: students) may be accessed through one of the following URLs.

- *../resources*: Used for Bulk operations (multi-object change requests)
- *../resources/12345*: Used for single object update or delete of Resource with ID 12345
- *../resources/Resource*: Used for single object create of a Resource (at which point the ID is unknown)

The Zone and Context of a Service are indicated in matrix parameters that appear only in the last segment of the Resource URL.

<sup>10</sup> The definitive sources for the per field mapping of all non-payload information to the URL and HTTP Header fields are Appendix C of the Infrastructure Services document and the SIF REST Developer Sandbox (SIF-RS).

### 3. SIF 3.0 Infrastructure Documents and Artifacts

After the basic overview of the SIF 3.0 infrastructure provided above, it is now possible to present a meaningful list of the various deliverables comprising the SIF 3.0 Infrastructure release.

#### 3.1. How is the SIF 3.0 Infrastructure specification documented?

The SIF 3.0 Infrastructure is described in four “volumes”, the first of which (this one) is descriptive, and the last three “normative” in that they define both the mandatory and optional requirements for SIF3.0-compliant software components

A brief overview of each is volume is contained in the table below.

Infrastructure Volume	Description	Primary Audience
<b>Read This First</b>	The overview, introduction and guide to the other SIF 3.0 Infrastructure artifacts. It is what you are reading now.	Anyone interested in understanding the functionality of the SIF 3.0 Infrastructure
<b>Base Architecture</b>	Defines the “core” concepts and detailed service operation framework of the SIF 3.0 infrastructure, and is the base document on which the other two infrastructure volumes identified below depend.	Those additionally needing to learn about the SIF 3.0 Infrastructure at a conceptual level.
<b>Infrastructure Services</b>	<p>Defines the complete specification (data structures, operations, and actions) for the set of directly accessible infrastructure services which together comprise the SIF 3.0 Environments Provider Interface (comparable to the SIF 2.x ZIS interface).</p> <p>This volume details which operations must be supported and which are optional for both the Direct and Brokered Environments. The most “important Infrastructure Services are:</p> <ul style="list-style-type: none"> <li>• <b>Environments:</b> Authenticates Consumer and defines Consumer Provisioning</li> <li>• <b>Requests Connector:</b> Routes all Authorized Consumer Requests to appropriate Service Provider.</li> <li>• <b>Queue Instance:</b> Allows the Consumer to synchronously access asynchronous Responses and Events</li> </ul>	Those additionally interested in learning about the core of the SIF 3.0 Infrastructure at a detailed level, or who need a reference for one or more specific Infrastructure Services.
<b>Utilities</b>	<p>Defines the additional set of Services providing secondary infrastructure functionality, which are accessed identically to Object Services.</p> <p>When added to the Infrastructure Services, the set of Service</p>	Those interested in learning about the full functionality of the SIF 3.0 Infrastructure, or who need a reference for one or more specific Utility

	<p>descriptions define the complete SIF 3.x infrastructure. The most important Utility Services are:</p> <ul style="list-style-type: none"> <li>• <b>Alerts:</b> Used when creating an exception reporting on an unexpected error condition</li> <li>• <b>Zones Registry:</b> Collection of all Zones available to the Consumer application..</li> <li>• <b>Providers Registry:</b> Collection of all approved Service Providers available to the Consumer application. Use of this Service is required for all self-provisioning Consumers..</li> </ul>	Services.
--	--	-----------

### 3.2. What purpose do the Infrastructure Object XML Schemas serve?

No surprises here. Like any SIF schema, these explicitly define the format of objects comprising the data payloads of messages exchanged when a Consumer is making requests of a Service. In this case however the Service is either an Infrastructure or Utility Service, and the payload relates to the Infrastructure rather than objects described in the Data Model.

Also as with any XML schema, when provided to an XML programmer toolset they result in the generation of marshaling and unmarshaling stubs that will automatically convert between the XML in a message payload and an “object” as defined in the programming language, and (where validation is enabled) optionally validate that the XML was / is in the correct format.

### 3.3. What purpose does SIF-RS (the SIF Rest Developer Sandbox) serve?

Documents and specifications can only go so far. The SIF 3.0 Developer REST Sandbox (SIF-RS)<sup>11</sup> started as an independently developed project based upon Java / Swagger technology and was specifically designed to lower the barriers to developing new SIF technology solutions in the following ways:

#### 3.3.1. Learning Tool

SIF-RS provides a comprehensive “hands on” learning tool for SIF developers, featuring a powerful GUI which “teaches by using”. The GUI functions as a Consumer, and communicates with the Direct Environments Provider component of the Sandbox.

A user can manually generate, view and alter:

- Valid XML instances of key Infrastructure, Utility and Data Model objects
- The associated URL segments, URL matrix parameters and HTTP header elements used in the construction and deciphering of SIF compliant messages.

<sup>11</sup> Probably the quickest way to come up to speed on the SIF 3.0 Infrastructure is, after reading this section, go to the SIF-RS URL (<http://rest3api.sifassociation.org/jsp/index.jsp>) and just “explore” with the provided GUI.



### 3.3.2. REST Platform Mapping Document

A great deal of effort was and will continue to be expended in ensuring that the XML instances and associated URL and HTTP header usage and formats made visible through the Sandbox GUI are valid and in conformance with the Infrastructure documentation.

As a result what the Sandbox user manipulates can be considered normative for the SIF 3.0 infrastructure specification, allowing SIF-RS to additionally function as the definitive “hands on” REST platform mapping specification for the SIF infrastructure. All SIF 3.0 compliant software implementations **must** support the XML payload and HTTP header fields and URL extensions in the manner illustrated in SIF-RS.

### 3.3.3. SIF 3.0 Application Testing Aid

The GUI may be replaced by a vendor supplied Consumer, allowing SIF-RS to serve as an Environments Provider against which Consumer applications can be debugged. SIF 3.0 Consumer applications can connect to the Direct Environment part of the Sandbox and attempt to register, issue Alerts, read and update object data and subscribe to and asynchronously receive object data change Events.

A limited debugging capability for both developer-created Direct Environments Providers (via the Sandbox GUI) and Service Providers (via extensions to the Sandbox Environments Provider) are planned.

## 3.4. What purpose will the SIF 3.0 Test Harness serve?

As noted, SIF-RS offers (among other things) a powerful debugging tool to allow developers to debug their SIF infrastructure code. Every Request, Response and Event message going over the wire is traced and can be compared to what the application being debugged is sending, correct (and explanatory) Error objects are issued when needed, key Infrastructure Services like the Queue Manager are implemented and available to service Consumer Requests, etc.

The SIF Test Harness compliments that by providing support for the subsequent phase of testing: the one between “can my software do anything?” and “does my software do everything?” It provides SIF developers with customizable test coverage for the specified range of functionality they believe their application supports, and it forms the basis of the SIF 3.0 Certification Test Suite.

This allows Developers to focus on the Test Harness when trying to detect problems with their software. When a problem is detected, the Sandbox should be used to isolate the cause and debug the proposed solution.

## 3.5. What other resources are available for developers?

There is a growing list of SIF 3.0 Infrastructure collateral which should be of use to developers. These currently include:

Resource	Location
A “Member Resources” area on the SIF website has been created specifically to provide additional resources, documentation and materials.	<a href="https://www.sifassociation.org/Resources/Member-Resources/Pages/default.aspx">https://www.sifassociation.org/Resources/Member-Resources/Pages/default.aspx</a>
Of specific interest are such topics as:	
- <b>SIF 3.0 Support</b>	
○ Common Terms	<a href="https://www.sifassociation.org/Resources/Member-Resources/SIF-Support/SIF_3-0_Support/Pages/default.aspx">https://www.sifassociation.org/Resources/Member-Resources/SIF-Support/SIF_3-0_Support/Pages/default.aspx</a>



<ul style="list-style-type: none"> <li>○ Core Differences between SIF 3.0 and SIF 2.x</li> <li>○ FAQs</li> <li>○ 3.0 Resources</li> <li>○ SIF 3.0 and JSON</li> <li>○ SIF Environments</li> <li>○ SIF REST – URL Structure</li> <li>○ Tools</li> </ul> <p>- <b>SIF-RS Tools and Reference Zone</b></p> <ul style="list-style-type: none"> <li>○ SIF 3.0 Concepts</li> <li>○ SIF 3.0 &amp; Standards</li> <li>○ Code Snippets</li> </ul> <p>These provide among other things, a quick start guide and additional useful references for developers utilizing SIF-RS.</p>	<p><a href="https://www.sifassociation.org/Resources/Member-Resources/SIF-Support/SIF-RS_Tools_and_Reference_Zone/Pages/default.aspx">https://www.sifassociation.org/Resources/Member-Resources/SIF-Support/SIF-RS_Tools_and_Reference_Zone/Pages/default.aspx</a></p>
<p>A set of AU community pages which contains examples and code snippets also of use to developers utilizing SIF-RS</p> <p>A set of additional AU developer resources</p>	<p><a href="http://sif.dd.com.au/">http://sif.dd.com.au/</a></p> <p><a href="http://kb.nsip.edu.au/display/SATWVC/Resources">http://kb.nsip.edu.au/display/SATWVC/Resources</a></p>
<p>Additional SIF 3.0 infrastructure collateral (presentations, design documentation) available to SIF members only</p>	<p><a href="https://www.sifassociation.org/Community/Global/Infrastructure/Pages/Shared-Documents.aspx">https://www.sifassociation.org/Community/Global/Infrastructure/Pages/Shared-Documents.aspx</a></p>

## 4. Aspects

This section examines several important aspects of messaging infrastructure in general (modularity, security, reliability and scalability). It highlights the multiple ways in which each aspect is supported by the SIF 3.0 Infrastructure, and provides references to the release documentation which describes the aspect in further detail.

Those readers desiring an understanding of one (or more) of these aspects as it relates to the SIF 3.0 Infrastructure should read the relevant subsection(s) below.

### 4.1. How modular are SIF 3.0 Environments?

A Consumer of a Service can be written to support only the functionality it needs. The situation is more complex for Environments Service Providers because the developers must correctly anticipate the level of functionality deemed critical to the set of applications registering themselves in the provided Environment. If no guidelines were specified, each Environments Service implementation could support a different combination of features, significantly reducing “out of the box” interoperability between SIF-conformant applications.

The SIF 3.0 Infrastructure standard addresses this problem by defining a small set of “Environment Service functionality packages” of increasing breadth and power. Each package consists of a set of required functionality that an Environments Provider could certify to, although a given product could optionally provide additional functionality as well.

For example, the requirements placed on an Environments Provider at a given SIF 3.0 site could range between supporting:

- A Direct Environment implemented by an SIS which services multiple instances of the same REST-based Dashboard Consumer application, where each instance runs on a different remote tablet device and each issues only a series of interactive Page *query* requests for Student data.
- A state-wide Brokered Environment implemented by extensions to an ESB middleware, product, which connects multiple regional and district SIS application Consumers dynamically issuing CEDS-compliant *create* and *update* Student data requests to the organization’s SLDS, and optionally receiving published Events describing all data changes.

In the SIF 3.0 architecture, the service functionality required of the first Environment is a completely contained subset of the service functionality required of the second. As a result, developers of the dashboard application can focus only on the manageable subset of the infrastructure documentation which relates to their specific needs, and yet still have their application work identically in both types of Environments.

#### **4.1.1. Environment Service “Functionality Levels”**

There are presently 3 levels of required Environment Services functionality defined, as highlighted in the table below. Planning for a SIF 3.0-based solution should start by determining which level is required to support the organization’s needs.

Environment Required Functionality	Low Mandatory (Direct Only)	Medium Mandatory (Direct or Brokered)	High Mandatory (Brokered only)
Change Operations	CRUD, Immediate IO Single object requests only	Add Multi-object requests, Events	Add delayed I/O
Query	By ID or all objects	Add Paged Query, predefined XQuery Template	Add dynamic XQuery Template definitions
Registries	None	Some exist, with prefixed unchanging contents (ex: Zone, Provider)	Most exist (ex: Namespace, External Code Lists), some with ability for Consumer to update (ex: Provider)
Queues	None	Queues with Long Polling support	Add Concurrent capability
Support for external Service Providers	None	Yes via support for “create” Provider Registry operation	Yes, and Administrative interfaces to manage external services
Security	HTTPS and Basic Authentication	Add HMAC SHA 256 Authentication	Add XML filtering

## 4.2. How Secure are SIF 3.0 solutions?

As in previous releases, SIF 3.0 was designed with Data Security baked in rather than being added on at the end. The Brokered Environment is essentially a “Data Confederacy” where data owners maintain control of their data and must “opt in” to share their data with other entities (application or organizational) rather than a “Data Union” with a single component that every application must share its data with.

At any site where SIF 3.0 is deployed, it is therefore the local data administrators who determine which applications (local and remote) can access or update sensitive data such as student discipline or health information, and exactly which parts of that data will be made visible.

The following table highlights some of the major security functionality provided / supported by the SIF 3.0 Infrastructure, and indicates where it is described in the specification.

Security Functionality	Description	Reference
<b>Authentication</b>	<p>The way in which the recipient of a message confirms the identity of the sender.</p> <p>In SIF 3.0, this occurs when an application initially registers as a Consumer with the Environments Provider, or when a Consumer issues a Request or a Service publishes an Event to the appropriate Connector.</p> <p>A Consumer's identity is a combination of Application, Application Instance and (where provided) User information. Multiple authorization methods involving a "shared secret" are defined.</p>	<p>The <i>Environments Service</i> section in Infrastructure Services.</p>
<b>Authorization</b>	<p>The way in which a Service Provider, having authenticated the identity of the Consumer issuing a Request, determines what access rights that Consumer has been given, and the ways in which limitations on those rights can be enforced.</p> <p>Site administrators can fine tune which Object Services a specific Consumer can access, and what types of access will be allowed</p>	<p>The <i>Environments Service</i> and <i>Provision Request Service</i> sections of Infrastructure Services.</p>
<b>Encryption</b>	<p>The way in which the contents of a message being routed over an insecure network is protected against being intercepted and understood by a 3<sup>rd</sup> party.</p>	<p>Enabled but not standardized. As stated:</p> <p><i>The Provider should use modern libraries and frameworks in order to support the likely high levels of encryption keys and certificate trust scrutiny; however the details of such support are no longer conveyed within the SIF standard</i></p>
<b>XML Filtering</b>	<p>The ability of the Environment Provider to selectively remove elements contained in a retrieved object before returning it to the requesting Consumer.</p> <p>This functionality has been an important part of enforcing site security policies at SIF 2.x deployment sites.</p>	<p>As with SIF 2.x, this functionality is enabled but the exact mechanism used is left unspecified. As stated:</p> <p><i>Detecting where to insert XML filtering, and determining what filter to apply are considered internal implementation details of required Environments Provider functionality. This logic may be divided among either the Connector or the Queue Infrastructure Service, but is not</i></p>

		<i>further described in this specification</i>
<b>Content Accountability</b>	Determining responsibility for data changes. In SIF 3.0 every element in a message payload can be traced back to the issuing application, and optionally, to the actual person responsible for the current value.	The <i>generatorId</i> element as described in the Base Architecture.

### 4.3. How reliable are SIF 3.0 Solutions?

Overall system reliability is of course dependent on the quality of the set of applications deployed in a given solution, and the extent to which these applications successfully interact.

The SIF Certification program provides the primary way to ensure seamless interoperability between SIF-certified components.

In addition, the SIF 3.0 architecture defines and standardizes a framework for constructing reliable multi-application solutions. The table below highlights the main components of this framework.

Reliability Framework Component	Description	Reference
<b>Guaranteed Message Delivery</b>	<p>The SIF standard requires that any issued message (Request, Response or Event) arrive at its destination, or (in unrecoverable error situations) that appropriate notification of the problem be given.</p> <p>In a Brokered Environment, this means a Service Provider can publish an Event, and it will eventually be delivered to every authorized subscriber, even if the recipient was off line or unavailable when the Event was posted.</p>	The <i>Request / Response / Event Message Exchange Choreography</i> section in the Base Architecture.
<b>Error Handling Analysis</b>	This analysis standardizes the process description of exactly how to handle errors detected in any part of the message exchange process, for Requests, Responses and Events.	The <i>Request / Response / Event Message Exchange Choreography</i> and the <i>Error Handling</i> sections in the Base Architecture.
<b>Alerts Service</b>	<p>The <i>Alerts</i> Utility Service provides a way for Consumers and Service providers to create an Alert object describing an exception condition which has just occurred. This could range from receiving an Event payload which failed to validate, to receiving an <i>update</i> request for an object that does not exist.</p> <p>The implementation of <i>Alerts</i> may range from appending</p>	The <i>Alerts Service</i> section in the Utility Services document.

	each created Alert to an Error Log file for later examination by site management, to converting and forwarding the Alert to a Preventative Maintenance and / or Service Management system.	
<b>Error Object</b>	<p>An Error object may be contained in place of one or more object results in the Response to an erroneous Consumer Request..</p> <p>The value of the internal Error code element matches one of the set of HTTP error codes utilized by REST programs, and the remainder of the Error object expands on the reason.</p>	The <i>Error Handling</i> section in the <i>Base Architecture</i> .

#### 4.4. How Scalable are SIF 3.0 deployments?

SIF 2.x Zone-based solutions have been deployed in a wide variety of educational organizations ranging from a single school to an entire State, for a wide variety of purposes. Performance limits were encountered at several particularly large deployments, especially at the end of a reporting period when message traffic volumes soared.

A careful analysis was made of the architectural causes of these limits, and all of the identified factors were addressed. There is excellent reason to believe that performance will be substantially improved for SIF 3.0.

The following table highlights some of the major performance enhancements contained in the SIF 3.0 Infrastructure, and indicates where they are described in the specification.

Performance Enhancement	Description	Reference
<b>Multiple Consumer Sessions for same application</b>	Site Administrators can enable the same application instance, using multiple <i>identifier</i> values (ex: <i>Johnson JHS</i> and <i>Miller Elementary</i> ) to register itself as multiple completely independent and parallel Consumers.	The <i>Environments Service</i> “create Environment” section in Infrastructure Services.
<b>Multiple FIFO Message Queues for same Consumer</b>	Any Consumer can set up separate Message Queues and assign different Event types to them.  This allows the less important Events (ex: Attendance) to sit in Queue B until the messages in Queue A have all been retrieved and processed.	The <i>Environment Scalability</i> section in Infrastructure Services.
<b>Multiple concurrent Consumer connections option to same Queue</b>	A Consumer may use (where the Queue Service supports it) more than one multiple simultaneous concurrent connection by having multiple HTTP Requests open with a single Queue.	The <i>Queue Service</i> section in Infrastructure Services.

<b>“Get Next and Pop” Queue Service Query Request</b>	This replaces the SIF 2.x “double handshake” between Consumer and Polling Queue with a single handshake, reducing the number of message exchanges necessary to retrieve an incoming Event during periods of high message traffic by a factor of two.	The <i>Queue Service</i> section in Infrastructure Services.
<b>“Long Polling”</b>	This option eliminates the message retrieval latency time previously associated with repeated polling for an arriving message, in periods of light traffic.	The <i>Queue Service</i> section in Infrastructure Services.
<b>Multiple Objects conveyed in single Request and single Event</b>	The ability to “pack” a single Request to or Event with multiple affected objects from a specific Service Provider offers huge performance benefits for several important use cases (ex: <i>a thousand small Attendance Objects created at the end of a Reporting Period result in the generation of 10 multi-object Events rather than 1000 single object Events</i> ).	The <i>Service Operation</i> section in the Base Architecture.
<b>Immediate Service Responses</b>	<p>In another major change from SIF 2.x, a Consumer may now request that the Service Response to any given Request be returned synchronously (i.e. immediately) on the HTTP Response to its issued HTTP Request, rather than appear asynchronously at some later time, intermixed with arriving Events.</p> <p>This will likely be the common case for most Requests, it conforms to the standard REST design pattern, it maximizes the effectiveness of Direct Environments and it is expected to improve overall Request / Response processing throughput significantly.</p>	The <i>Basic Infrastructure Framework</i> section in the Base Architecture.
<b>eTag</b>	<p>An eTag is equivalent to a “checksum” on <u>all</u> the objects of the type being queried which are maintained by the Service Provider, and it is optionally returned in all Query Responses.</p> <p>If a valid eTag is included in a Query Request, and if there were no object data changes since the eTag was created, no objects are returned.</p> <p>This offers an efficient way for Consumer applications executing on mobile devices (which are unlikely to subscribe to Events) to determine whether they already have the latest data on one or more objects without comparing all the elements.</p>	The <i>Service Operation (Query)</i> section in the Base Architecture.

## 5. Migration

The SIF 3.0 release offers adopters significant enhancements to the functionality contained in prior releases of the infrastructure. However it also represents the first time in almost a decade that the SIF Association will be packaging new infrastructure functionality in a form that is not backwards compatible with earlier versions.

Any reader with a currently deployed SIF v2.x solution or SIF v2.x conformant application should read this section.

### 5.1. Are there a set of guidelines available to end user seeking to migrate to the new release?

While the development of the SIF 3.0 Migration Guidelines document is still a work in progress, some clear migration rules and techniques have been developed by the SIF Migration Project Team. These are being used in various Proof of Concept (POC) projects, and should be of use to early technology adopters as well.

Any SIF 3.0 Infrastructure Migration strategy must be:

- **Seamless.** Existing SIF 2.x Zone functionality must not break
- **Incremental.** Only one Zone component need be upgraded at a time
- **Justifiable.** A clear value proposition must exist for each step in the migration process
- **Practical.** A real world End-User centric SIF migration use case must be addressed

In addition, the following represent some valuable “lessons learned” in terms of constructing a migration strategy.

#### 1. Some 2.x Zones may never upgrade to Zones in SIF 3.x Environments

Where a SIF solution is deployed and working and there is no compelling reason to upgrade existing functionality, there is no compelling reason to migrate. Additionally, some SIF 2.x objects may never be part of the CEDS Logical Data Model (ex: Food Service), providing vendors of SIF-certified products in those spaces with little motivation to re-code for, and then re-certify to, the new release.

#### 2. Version “Bridging” can be achieved in multiple ways.

**Middleware Broker:** This is the traditional approach to bridging SIF applications supporting different releases. It puts all the conversion logic into the middleware broker, where it executes as each message is routed. Basically the Zone spans releases but applications do not.

**Application Data Store:** A viable alternative is to have the application register itself in both a SIF 2.x Zone and a SIF 3.0 Environment. Bridging is then done in the application’s Data Store and not in the middleware. The Application spans releases but the Zone does not.

### 5.2. Is there an example of a SIF 3.0 Infrastructure Migration Strategy?


The following 4-step example represents a “bottom up” approach to SIF 3.0 Infrastructure migration. There are several “universal properties” embedded in this strategy that appear relevant to other SIF migration use cases as well.



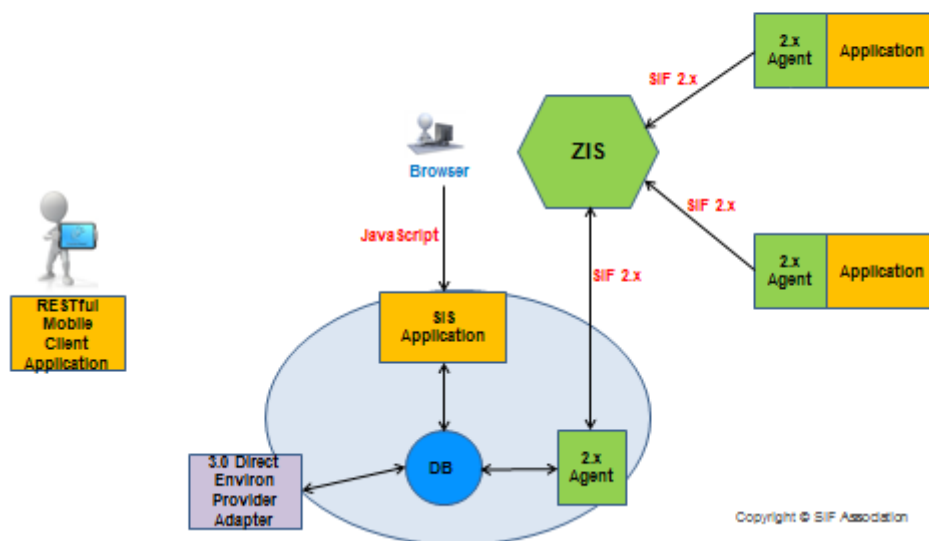
- Migration starts with one SIF 2.x Object Provider at a time
- The Object Provider is relatively unaware of its SIF 2.x Agent. They are bound together primarily through the Agent manipulation of the Application's Data Store.
- A SIF 3.x Adapter is added in a similar manner such that the main component binding the Adapter, the Agent and the Application together is the Data Store.
- Multiple SIF 3.0 Direct Environments will appear first, and only later be merged into a SIF 3.0 Brokered Environment.
- The conversion from a Direct Environment to a Brokered Environment is transparent to all Consumers of the Direct Environment.

The sequence of Migration milestones in this particular Migration Strategy are shown below:

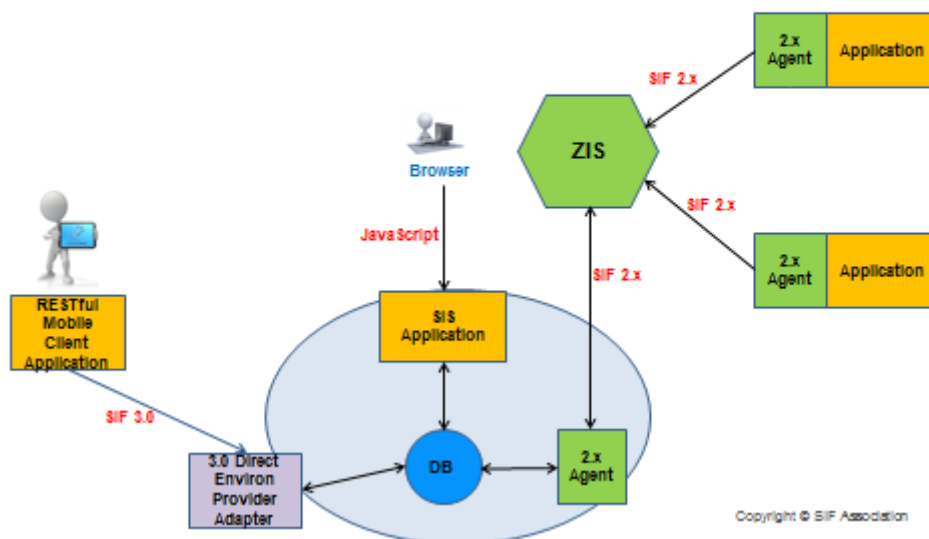
### Migration Motivation: An SIS must support RESTful Clients executing on Mobile Devices



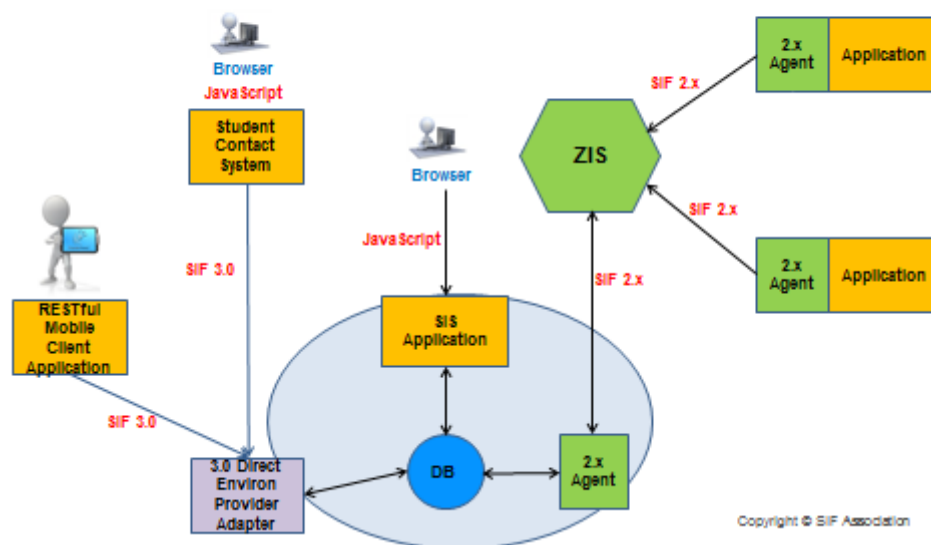
## Step 1: Add a SIF v3.0 Direct Environment Provider Adapter



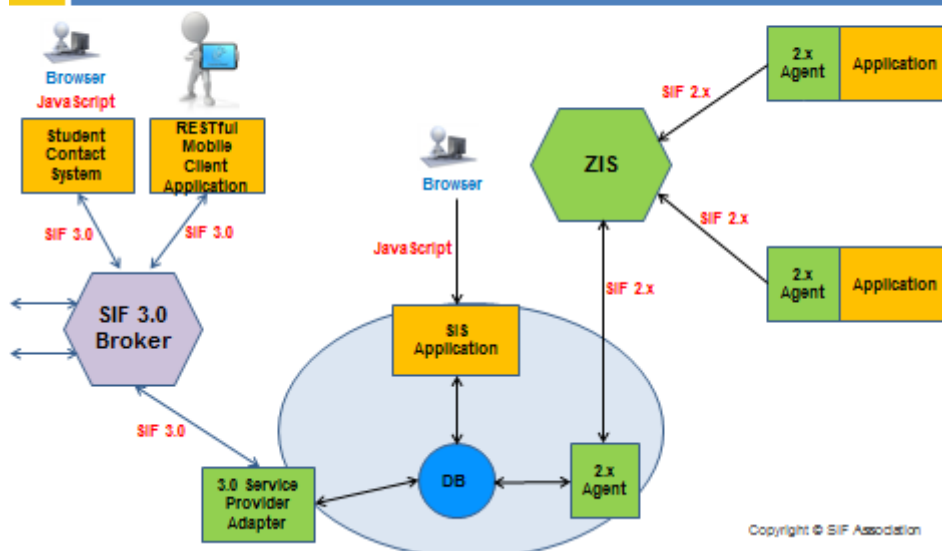
## Step 2: Add SIF-Compliant RESTful Resource Requests to Mobile Client Applications



### Step 3: Deploy other RESTful Client Apps to access and update SIS data directly



### Step 4: Upgrade to SIF Brokered Environment and deploy other RESTful Client Apps



At this point, a Student Object *create* request from the SIF 3.0 mobile client application will:

- Be dispatched through the SIF 3.0 Brokered Environment to the SIF 3.0 Service Provider Adapter.
- The adapter will update the SIS database

- The database update will in turn “trigger” the SIF 2.X Agent to publish a separate Event for each object successfully created, to all subscribers in the SIF 2.x Zone
- The SIF 3.0 Service Provider Adapter will publish the single multi-object Event resulting from the change it made to the internal SIS database, to all subscribers in the SIF 3.0 Environment.

Note that in this scenario, there is no involved on-the-wire conversion, and all SIF 3.0 upgrades are totally transparent to the remaining older components. In the final step, which may not happen for years (if ever), the last SIF 2.x application is upgraded, everything hooks in through the SIF 3.0 Environment Broker, and the SIF 2.x ZIS can be safely retired ... seamlessly and incrementally.

## Appendix A: Miscellaneous Questions

The following questions did not fit into any specific technical section, but the answers were felt to be of general interest.

### A1. Why was there such a delay between major SIF Releases?

The SIF 3.0 Infrastructure specification represents the first new “major version” release of the SIF infrastructure since 2007, although new minor releases of the SIF 2.x specification have traditionally been issued every 9-12 months. There were two main reasons for the delay.

#### The SIF Standard has a large deployed base

At times *“the present can be the enemy of the future”*. The fact that the existing SIF Community is composed of members who were and are generally satisfied with the functionality provided by the SIF 2.x release infrastructure, meant there needed to be a compelling reason to consider upgrading their working solutions and requesting the development of new products to support a potentially non-backwardly compatible major SIF infrastructure release.

With the advent of CEDS, that compelling reason appeared for the SIF US community. It complimented the growing need in all locales (US, AU and UK) for a secure and scalable REST based infrastructure that would “capture” applications that would otherwise adopt one or more low-end alternatives, so that they could be utilized in future SIF-conformant solutions.

#### The SIF Association is an Open Standards Organization and SIF is an Open Standard

Development of a truly Open Standard involves the collaborative efforts of interested developers from a wide variety of end user and vendor organizations, all focused on jointly working out and reaching consensus on the details of the functionality of the next release. When completed, such an open standard creates a level playing field for vendor adopters and allows end users to select “best of breed” conformant products for their deployed solutions.

In contrast, alternative vendor-specific “open source” projects and foundation-specific initiatives focused around specific product implementations have an easier time of it. They generally seek community feedback but only a small internal group is responsible for:

- All technical decisions
- All the content contained in the standard documentation
- Retaining the resulting IP (everyone outside that group is often asked to sign an IP agreement).

Removal of the need to seek consensus before defining and documenting the specification results in a more rapid turnaround than can be achieved by the open standards approach. However it usually locks in one or more software components in every solution, inhibits competition, and removes the end user’s ability to select “best of breed” replacements for those core modules whose interface is being standardized.

So it took us longer. But we’re here now.

### A.2 Isn’t SIF 3.0 an overly complex infrastructure specification?

The documentation for all three volumes of the SIF 3.0 Infrastructure specification totals around 200 pages. The primary reason for this was the need to standardize how the extensive level of functionality required to support large State Level educational solutions would be provided.

## Packages within Packages

Sometimes it takes a lot of careful design to make something simple.

SIF 3.0 functionality has been specified as a modular collection of RESTful service interfaces, packaged together in one of three layers of increasing power, each building on the last. Since each functionality layer is a completely contained subset of the one “higher up” in the functionality strata, application developers need to understand and utilize only that subset of the infrastructure documentation which relates directly to their needs. Their resulting applications will work identically in any deployment supporting at least the set of functionality in the package they are designed to, or any level above that.

### Minimal Package

The “lowest level package” on which all the others build is very lean indeed. It defines the minimal set of functionality that both a SIF-conformant client and service are each required to provide in order for them to be considered “SIF 3.0 Infrastructure compliant”, and interoperate successfully. An example of such a “minimal infrastructure” SIF-conformant client might be a student portal application deployed on a mobile device, used in place of a Browser GUI.

To implement the bottom layer, developers new to SIF can read this RTF, check out the SIF-RS Sandbox for a hands on example of exactly what is going over the wire, and apply their REST knowledge to start developing a SIF-compliant solution almost immediately.

### Leveraging Infrastructure Vendor Offerings

It is also anticipated that in a manner similar to how earlier releases evolved, much of the SIF 3.0-specific infrastructure support will likely be provided by dedicated SIF infrastructure product vendors or open source SIF-adaptor projects. This will leave many educational application developers free to focus primarily on the functionality of their product and how it bridges to the relevant elements of the SIF Data Model version which they support, rather than on the details of the infrastructure over which their payloads travel.

## A.3 Why was a SIF 3.0.1 release needed?

After a major release like SIF 3.0 is issued, the experiences of early adopters often indicate that parts of the specification cannot be implemented easily on important developer tool sets, or that some important components of functionality were inadvertently omitted.

The solution has historically been to issue a “fix release” to correct these problems, before working on any minor release extensions. Perhaps the best way to answer this question for SIF 3.0.1 is to refer to section 1.13 of the Base Architecture Guide, which contains the list of the important fixes and new functionality which it introduced.