



SIF 3.0.1 Infrastructure: Utilities

Table of Contents

1.	Introduction	4
1.1.	Required Reading.....	4
1.2	Utility Service Types	4
1.2.1	Administrative Level Authorization.....	5
1.2.2	Scope	5
1.2.3	Functionality	6
2.	Zones Registry Utility Service (zonesRegistry)	8
2.1	Zones and Consumer Requests	8
2.2	Presence and Scope	8
2.3	Supported Operations	9
2.4	Zones Registry Data Elements	9
2.5	XML Example	9
3.	Providers Registry.....	11
3.1	Supported Operations	12
3.2	Providers Registry Data Elements.....	12
3.3	XML Examples.....	15
4.	Namespaces Registry	17
4.1	Service Implementation Strategy	17
4.2	XML Schema Snippet.....	17
4.3	Supported Operations	18
4.4	Namespaces Registry Data Elements.....	18
4.5	XML Examples.....	19
5.	Code Sets Registry.....	20
5.1	Service Implementation Strategy	20
5.2	Alternatives	20
5.3	Supported Operations	21
5.4	Code Sets Registry Data Elements	21
5.5	XML Example	24
6.	XQuery Templates Registry	26
6.1	Security	26
6.1.1	Lazy Authorization	27
6.1.2	Security Limitations.....	27
6.2	Restrictions on Use	29

Singular Form.....	29
Formula Form.....	31
Extended Form.....	31
Parameters.....	31
Access and Ownership.....	31
6.3 Supported Operations	31
6.4 XQuery Templates Registry Data Elements	32
Example	33
7. Alerts	34
7.1 Supported Operations	34
7.2 Error Handling.....	35
7.2.1 A Service Provider receives a Consumer operation invocation which it rejects	35
7.2.2 The Consumer receives a Service Provider Response which it rejects.....	36
7.2.3 A Service Provider posts an Event which a Subscribing Consumer rejects	36
7.2.4 A Service Provider detects an inactive Consumer	37
7.3 Alert Object Data Elements.....	38
7.4 XML Examples.....	40
Appendix A: Minimal (Simple) Consumer	42
A.1 Minimal Consumer vs. more “demanding” Consumers	43
A.2 Core Infrastructure Service Interactions	43
A.3 Utility Service Interactions.....	45
Appendix B: Minimal Environments Providers	47
B.1 Overarching Design Constraints.....	47
B.2 Minimal Environment Architectural Requirements.....	48
B.3 Core Infrastructure Service Support	49
B.4 Utility Service Interactions.....	50
Appendix C: Minimal Service Providers	53

1. Introduction

A SIF 3.0 Utility Service is a Data Service that *provides* a data object related to the SIF infrastructure, rather than an object which is part of a particular locale-specific data model.

- They support some or all of the identical set of Consumer Requests (Query, Create, Update and Delete).
- Some Utility services may be subscribed to, and publish object Events when their internal data changes
- They are accessed by Consumers through the *requestsConnector* and return any delayed Responses and Events into the specified *Consumer Polling Queue*.
- They can be located by Consumers through a *Provider* entry in the Global Zone section of the *Service Providers Registry*.
- They are documented with the infrastructure rather than the set of educational object services which use that infrastructure, and are present in every release based upon that infrastructure version, independent of locale.
- They may be implemented independently of the core Infrastructure Services or any other Utility Service, and integrated into a site-wide solution in the same manner as implementations of other data object Services.
- They have their formats defined within the Infrastructure namespace (<http://sifassociation.org/infrastructure/3.0>), rather than in one of the Data Model namespaces which define the payload of the other Object and Functional Services.

1.1. Required Reading

Please refer to the *Basic Architecture* document for an understanding of the terminology, concepts and global XML element definitions, Service types, operation descriptions and exchange choreographies that will be referred to here. Additionally, the legal notices contained in the Basic Architecture apply to this document.

Please refer to the *Infrastructure Services* document to gain an understanding of the “core” Infrastructure Service components, which these Utility Services supplement.

Some familiarity with both those documents is assumed as a prerequisite, as all Utility Services will employ the Service framework and leverage the specific Infrastructure Services described there.

1.2 Utility Service Types

As noted, all Utility Services are accessed through the *Requests Connector* Service and follow the standard Data Object Service API in that they support some or all of the defined Object Service Provider interface methods (Query, Create, Update and Delete). There are two additional characteristics which are unique to this set of Services.

1.2.1 **Administrative Level Authorization**

Particularly in Brokered Environments, there may be one or more administrative Consumer applications which are used to configure and maintain the Environments of other Consumers. These applications may be granted authorization rights to additional Provider Service CRUD Methods unavailable to non-administrative Consumers.

Where they exist, such extensions are site specific and except in a few cases, will not be explicitly defined in this document.

1.2.2 **Scope**

All Utility Services span the entire Consumer's Environment. They each have a single entry in the Service Providers Registry, located in the *environment-global* Zone.

This means that any in any Consumer Request, the combination of a *serviceType* value of "UTILITY" and the *serviceName*, provides enough information to allow the Connector in a Brokered Environment to first determine the identity of the Utility Service Provider which should receive it, and then construct the exact Service URL to route it to.

As a result, any *zoneld* provided by the Consumer in a Utility Service Request does not affect the ultimate destination of the request, and can be treated differently than if the *serviceType* value was "OBJECT" or "FUNCTIONAL". There, a Consumer-provided *zoneld* may result in the Connector forwarding the Consumer's Request to an entirely different Service Provider URL (ex: an SIS for Special Ed students or one spanning only a specific middle school).

When a Service Consumer issues a Request to a Utility Service Registry however:

- If the Consumer explicitly specified the *environment-global* Zone, a non-qualified Query is assumed and the Utility Service returns all entries in its Registry. If the Consumer specified any other legal Zone (or its default *zoneld* was supplied by the Connector) only those entries applicable to that *zoneld* are returned. Essentially the *zoneld* parameter serves as a query constraint on the entries in the Registry.
- For the Service Providers Registry, this means a Consumer gets only the Provider entries for its default Zone or, if it specified the *environment-global* Zone, it gets all Service Providers for all Zones.
- For the External Code List Registry (and others), things are somewhat different. The default assumption is that all External Code Lists are applicable to all Zones, and are therefore "Global" in scope, and their entry has a *zoneld* value of *environment-global*.

Whenever this is not true (for example grade enumerator codes may be different for different school Zones in the District Environment) a Zone-specific External Code List entry of the same name but with the different set of values is created, and tagged with the corresponding *zoneld* value.

When a Zone-specific query request is issued, each entry in the global list is then examined in turn

- If there is no corresponding Code List with that name defined specific to the selected Zone, the Global Code List entry is returned.
- If there is a specific entry for that Code List in the selected Zone, that entry is returned instead.

The end result is that when a Consumer issues a request to a Utility Service with a *zoneld* other than *environment-global*, only those entries (ex: Providers, External Code Lists) relevant to the specified Zone are returned. If the *zoneld* specified is *environment-global*, all entries across all the Zones in the Consumer Environment will be reported.

1.2.3 **Functionality**

The generic functionality provided by each Utility Service in supporting the Environments Provider API is shown in the following table. Unless otherwise noted, no Utility Service supports XQuery or Paged Query Requests.

They will all be described in further detail in the sections that follow.

Utility Registry Services	Functionality	Access
Zones	Contains the name and description of all Zones in the Consumer's Environment (not all of which may be accessible by the Consumer).	Only Query access is enabled for non-Administrative Consumers. Change Requests are not supported.
Providers	Contains entries with the name, service type, vendor information and any supported Context of all Service Providers in a given Zone (not all of which may be accessible by the Consumer).	Only Service Providers operating in a Brokered Environment may create and change the contents of their Registry Entries. All Consumers in both types of Environment have Query functionality.
Namespaces	Used to retrieve namespace URIs and their corresponding schema location URLs, currently valid within the zone.	Only Query is enabled for non-Administrative Consumers. Change Requests are not supported.
Code Sets	Validates specific code values claimed to be part of an identified named code set, used to reduce breaking changes introduced via normative references to external sources like the NCES Handbook.	Only Query is enabled for non-Administrative Consumers. Change Requests are not supported.

XQuery Templates	Accepts, validates (per-Consumer) and provides access to previously submitted or pre-defined XQuery Templates.	All approved Service Providers can use the supplied Token to retrieve the corresponding XQuery Script using Query by (XQuery Template Token) ID. This script is then used as the basis of expanding and executing a Consumer issued XQuery Request containing that Token and a set of values corresponding to the parameters in the template. Only Query is enabled for non-Administrative Consumers. Change Requests are not supported.
Other Utility Services	Functionality	Access
Alerts	Accepts and forwards reported alerts (notifications warnings and errors) in accordance with the pre-specified diagnostic and security policies defined by the site administrator.	All Service Consumers can create an Alert, and issue Query by (Alert) ID to access their own previously reported Alerts. Only administrative Consumers can access and change the Alerts issued by other Consumers.

2. Zones Registry Utility Service (zonesRegistry)

The Zones Registry Utility Service contains the name and description of all Zones visible to a registered Consumer within its Environment. Some Zones may contain services which are not accessible by the Consumer due to authorization restrictions.

2.1 Zones and Consumer Requests

A *Zone*¹ is basically a collection of Service Providers and associated Utility Registry information within the Consumer's Environment, pre-organized by the site Administrator to correspond to a discrete hierarchical grouping within the owning educational organization such as a school or district, or an alternative grouping such as the set of applications supporting Special Ed students. Zone identifiers are chosen by the administrator and can follow any convention that best meets the needs of the deploying organization.

Each Data Object and Functional Service "instance" accessible within the Consumer's Environment is scoped to a Zone, although a given Service Provider implementation may support the same Service Provider interface in several Zones. As noted above, there is one "special" Zone (*environment-global*) which is reserved for "globally available" (i.e. Environment-wide) Utility Services.

Every Consumer request for any Provider Service is always qualified by the Zone in which the Service is to be found. Each Service Consumer is assigned a "default" Zone at Registration time, which is used whenever a specific Zone is not explicitly included in one of its Provider Service Requests. If there is no matching registered Service Provider for any Consumer Request, the request must fail.

2.2 Presence and Scope

The presence of a Zones Registry Utility Service is mandatory for all Environments that support Consumer self-provisioning, whether Direct or Brokered. In the simplest case (typically a Direct Environment provided by an SIS or LMS application), the Zones Registry might consist solely of the Zone ID and Zone Description (with no additional parameters) for two Zones:

- *environment-global*
- XXX

In this case "XXX" would be the "default" Zone assigned to all registered Consumers. Its value might represent the name of the application providing the non-utility services (such as Assessment).

A Consumer does not need to access the Zones Registry if it interacts solely with the set of Service Providers contained in its assigned "default" Zone, and the available Utility Services in the *environment-global* Zone.

¹ Please refer to the SIF 3.0 Base Architecture document for a more complete description of a SIF Zone.

2.3 Supported Operations

The set of Zones is fixed for each Registered Consumer, and does not change while the Consumer is active. Therefore the Zones Registry Service is not required to support change requests or to publish change events.

The only guaranteed operation for non-Administrative Consumers is Bulk Query. There is no support for dynamic Query, and there are no defined Zones Registry Service Paths or XQuery Templates.

2.4 Zones Registry Data Elements

The following data elements comprise the contents of the Zone Service Registry.

Element	Char	Description	Type
zones		Top level element of the Zones Registry	Collection
zones/zone	RM	Individual entry defining a specific Zone within the Consumer's Environment	Zones Registry Entry
zones/zone@id	M	The unique id of the Zone, and the key to the Zone element in the Zones Registry. Typically this is the globally unique and recognizable "name" of the Zone, rather than a meaningless code.	xs:token
zones/zone/description	O	A short (possibly multi-sentence) description of the Zone	xs:normalizedString
zones/zone/properties	O	An optional set of free-form name value properties further defining the Zone element	
zones/zone/properties/property	MR	The value of a specific property which defines the Zone.	xs:token
zones/zone/properties/property@name	M	The name of that property for which the value applies	xs:token

2.5 XML Example

The payload of the Response to Query Request issued to the Zones Registry Service in a minimal Direct Environment is shown below.

```
<zones>
  <zone id="environment-global">
    <description>Contains all global Utility services</description>
    <properties>
      <property name="type">Utilities</property>
    </properties>
  </zone>
</zones>
```

```
</zone>

<zone id="SuffolkMiddleSchoolSIS">
  <description>Scopes SIS Service Provider at School</description>
  <properties>
    <property name="Special Ed">Student</property>
    <property name="administrator">alexj@ SuffolkMiddleSchool.va.edu</property>
  </properties>
</zone>
</zones>
```

3. Providers Registry

There is one *Providers Registry Utility* Service per Consumer Environment. The Registry contains a list of Provider Entries, each of which includes:

- A Zone ID
- The type of the provided Service in that Zone (Utility, Data Object, Functional, Service Path or XQuery Template).
- The URL identification of the provided Service.)
- Any specific Context that Service supports.
- Any Service functionality extensions (such as support for Dynamic Queries or the ability to provide “Total Count”)

All potentially accessible Services have an entry in the Providers Registry (including the Providers Registry Utility Service itself), although full or even partial Consumer access to that Service is determined by the access rights currently granted in the Consumer’s Environment object², and is not guaranteed.

As noted in the Basic Architecture document sections on Service Types, there are 5 types of Services which could be represented in the Provider Service Registry.

Service Type	Defined	URL Service Identifier in Registry Entry (example)	Operations
Data Object	Data Model Schema	students	Query, Create, Update and Delete
Utility Object	Infrastructure 3.0 Schema	zones	Defined in this document (Query only in this case)
Functional Service	Data Model Schema and Functional Description	studentRecordExchanges	Specific to the Function
Service Path	Data Model Binding for Object Type. Schema identical to schema of object type name in final URL segment	<i>sections/{}/students</i> where “{}” indicates the section to report Students for	Read Only
XQuery Template	Data Model Binding defines Schema and XQuery script contents	StudentSnapshot	Read Only

² Please refer to the Environments Service section in the Infrastructure Services document.

3.1 Supported Operations

In many cases, a Service Consumer will be pre-provisioned to be able to access the set of Service Providers it must rely on to perform its functions. In that case, it does not need to utilize the Providers Registry Service.

However at sites where there the security policy is “authorization on demand”, any Service Consumer with the proper authorization rights may be expected to dynamically utilize the Providers Registry to discover available Service Providers in its default Zone or elsewhere. At that point it can dynamically self-Provision itself to issue Requests and (where applicable) to subscribe to Events from one or more of these Providers

Such usage requires the ability to Query the Providers Registry.

If the Consumer has the proper authorization (and is deployed in a Brokered Environment), it can also register itself as a Service Provider by creating one or more of its own Provider Entries that will in turn be visible to other Consumers³. If the Service Provider implementation supports multiple service options (multiple Zones, multiple Contexts, multiple Data Model version numbers, etc.) it **must** create multiple corresponding entries in the Provider Registry.

The Providers Registry Service must publish change Events when Service Entries are added, updated or deleted. The Providers Registry Service itself does not support dynamic Query, and there are no defined Providers Registry Service Paths or XQuery Templates.

Operation	Description
Query	Provides the Consumer with a list of all accessible Service Providers within the specified Zone, or (if access was through the environment-global Zone), a list of all Service Providers of all types in all Zones.
Create (Brokered Environments only)	Allows a Consumer to initially provision (or attempt to provision) itself as a Service Provider. As in any Create Request, multiple Provider Entries may be created with this call (as when the Consumer requests to be a Provider of a given object type for multiple Contexts). Success results in the creation of one or more new Provider Entries.
Delete (Brokered Environments only)	Removes one or more specified Entries in the Providers Registry. Except for Administrative Tools, this Request must be rejected for all Consumers other than the one which created the entry in the first place. Success results in the deletion of the specified entries from the Providers Registry. At this point these Services are unavailable to all Consumers ⁴ .
Update	Prohibited

3.2 Providers Registry Data Elements

The following data elements comprise the contents of the Service Providers Registry. As noted, the entries returned in the response to a Consumer Query are scoped by the Zone in which the Providers Registry is accessed.

³ This is very similar to when a SIF US 2.6 application provisions itself as an Object Provider.

⁴ Such an Entry Deletion request is generally issued as part of an ordered “unregister” sequence for the Provider, prior to its deleting the Environment. It is not required that a Provider issue this however.

Element	Char	Description	Type
providers		Top level element in the Providers Registry	Collection
providers/provider	RM	Individual entry defining a specific provider	Providers Registry Entry
providers/provider@id	M	The unique id of the Provider entry. This is typically only used by the Service Provider to directly Query or Delete its own entry It may be a hashed combination of the next 5 entries in this table or it may be a standard UUID.	xs:token
Service “Scoping” Elements			
providers/provider/serviceType	M	The “generic” type of Service being provided. (Note: the namespace of Utility objects is fixed by the infrastructure version. The namespace of all object and functional Services must match the one the Consumer gave when it registered, and are not repeated in the Providers Registry)	One of: UTILITY OBJECT FUNCTIONAL SERVICEPATH XQUERYTEMPLATE
providers/provider/serviceName	M	The part of the URL in the Consumer Request that, along with the Context and Zone must uniquely identify the recipient which will provide the response. This element is an XML Token, whose exact meaning is determined by the Service Type. In all cases this value determines exactly what the format of the message payloads will be. Examples: If the Service Type is: <ul style="list-style-type: none">• UTILITY: then zones is the Zones Registry Utility Service (one of a fixed list of services)• OBJECT: then sections is the Sections object provider• SERVICEPATH: then sections/{}/students will return all Student objects in the Section who’s RefId is supplied in place of the “{}” in the URL segment between <i>sections</i> and <i>students</i> (ex: <i>sections/1234/students</i> returns all Students in Section 1234) XQUERYTEMPLATE: then StudentSnapshot is an XQuery Template Token which is predefined to return a specific report.	xs:token

		FUNCTION: then StudentRecordExchange is a Functional Service which creates a unique job object for the issuing Consumer which encapsulates the behavior of what can be a multi-step process.	
providers/provider/contextId	M	The value of the Service Context being provided (if any). If there is no context associated with this object type, the value of this element is "DEFAULT".	xs:token
providers/provider/zoneId	M	The Zone in which the Service Context is being provided. This must correspond to one of the entries in the Zones Registry. The combination of ZoneId, Service Type, Service Name, Context Name and (if present) Context value must be unique for every entry in the Service Providers Registry. .	xs:token
"Functionality Description" Elements			
providers/provider/providerName	M	The name of the Service Provider as it would be referred to by the administrator of the Zone (ex: RamseySIS).	xs:token
providers/provider/querySupport	M	A set of elements corresponding to the extent of Service Provider support for Consumer Query Requests (above and beyond support for Query by Id, which is mandatory and assumed).	
providers/provider/querySupport/dynamicQuery	M	Present if this Service is capable of processing a Dynamic Query "where" argument when included in the URL of a Query Request. ⁵	xs:boolean
providers/provider/querySupport/paged	M	True if this service is capable of responding to Paged Query Requests.	xs:boolean
providers/provider/querySupport/maximumPageSize	MC	If Paged Queries are supported, this is the maximum number of Objects that will be returned on a Page of Query results.	xs:unsignedInt

⁵ The actual level of Dynamic Query support may vary in terms of whether all comparators are supported or only "=", whether multiple "and / or" Boolean expressions are supported, and how deeply nested the identified elements in the "where" argument can be. This level of detail is specific to the Object type and may be also be defined in the Data Model "binding".

providers/provider/querySupport/totalCount	MC	If Paged Queries are supported, this indicates whether the Service Provider will return a “total count” of all objects satisfying the original Query, either in each Paged Response, or the Response to a Paged Request with a Page Size of zero.	xs:boolean
“Product Description” Elements (these are taken from the Provider’s Environment and are omitted from the Provider Entry Create but returned in the Query)			
providers/provider/applicationProduct	C	Application Vendor Identification	productIdentityType ⁶
providers/provider/adapterProduct	O	Adapter Vendor Identification	productIdentityType
“Hidden” Elements (this is included in the Provider Entry create but not returned in the Query)			
providers/provider/endpoint	MC	In a Brokered Environment, this element contains the URL where all requests for this Service Provider should be re-invoked by the Connector (which received them from the Consumer). Note: this URL is independent of Zone and Context information, which is supplied elsewhere in this entry.	protocolType ⁷

3.3 XML Examples

1. The payload of a *create* Provider request issued to the Providers Registry Service is shown below. Note the end point specification.

```
<providers>
  <provider= “”>
    <serviceType>OBJECT</serviceType>
    <serviceName>students</serviceName>
    <contextId>DEFAULT”</contextId>
    <zoneId>SuffolkMiddleSchoolSIS</zoneId>
    <providerName>RamseySIS</providerName>
    <querySupport>
      <dynamicQuery>true</dynamicQuery>
      <totalCount>>false</totalCount>
      <paged>true</paged>
      <maxPageSize>3000</maxPageSize>
    </querySupport>
    <endPoint>https://www.suffolk-jhs.pedmontdistrict.nc.edu/suffolkmiddleschool/sis</endPoint>
  </provider>
</providers>
```

⁶ See Appendix D of the Infrastructure Service Document for an expansion of this Data Type.

⁷ See Appendix D of the Infrastructure Service Document for an expansion of this Data Type

2. The payload of the Response to a Query Request which corresponds to the RamseySIS entry created above. Note the endPoint element is not shown, and application identification information obtained from the RamseySIS Consumer Environment is included in this entry.

```
<providers>
  <provider id="ab1234cd5678ef90cba">
    <serviceType>OBJECT</serviceType>
    <serviceName>student</serviceName>
    <contextId>DEFAULT</contextId>
    <zoneId>SuffolkMiddleSchoolSIS</zoneId>
    <providerName>RamseySIS</providerName>
    <querySupport>
      <dynamicQuery>true</dynamicQuery>
      <totalCount>false</totalCount>
      <paged>true</paged>
      <pageSize>3000</pageSize>
    </querySupport>
    <applicationProduct>
      <vendorName>Raffles Software</vendorName>
      <productName>RafflesSIS</productName>
      <productVersion>1.5</productVersion>
    </applicationProduct>
    <adapterProduct>
      <vendorName>SifRus</vendorName>
      <productName>GlobalAdapt</productName>
      <productVersion>4.2</productVersion>
    </adapterProduct>
  </provider>
</providers>
```


4. Namespaces Registry

The Namespaces Registry contains the set of XML namespace URIs and their corresponding schema location URLs which are currently valid within the Environment. Both Service Consumers and Providers **SHOULD** stay synchronized with this registry in order to prevent and detect namespace mismatches while exchanging messages.

In terms of guaranteeing interoperability and enforcing element-level privacy restrictions, controlling the Namespaces contained in exchanged messages is as important to Environment Administrators as controlling the XML elements themselves, because Consumer to Provider interoperability requires standardization of the name, value and scope of every element exchanged.

This imposes a Consumer to Provider co-dependency, in that each has to construct messages containing only the namespaces that the partner has the schemas for and can validate. In general, such schema agreement is static and predefined by either the version of the Infrastructure or Data Model supported within the Environment. In this case, neither the Consumer nor Provider need access the Namespaces Registry.

However in Environments where namespaces are subject to change, where customized profile extensions to the core SIF schemas have been made, and especially where the Environments Provider, Data Warehouse or Service Provider is bridging between versions, the Namespaces Registry allows for both synchronization and just in time retrieval of needed schemas, to minimize rejection of messages containing unexpected namespaces.

In particular, with the separation of infrastructure and data model in SIF 3.0, there is no longer one Namespace that all Servers and Providers within the Environment must conform to. The Namespaces Registry is provided as the only sure way for a Consumer determining what the Providers within the Environment will expect to see. It allows Service Consumers wanting to place XML on the wire or retrieve data over the wire to dynamically determine whether the namespaces (including version) they rely on are supported within the Environment before they begin operation.

4.1 Service Implementation Strategy

Each entry in the Namespaces Registry carries only the URL of a single schema file, which by convention leverages the XML Schema Namespaces themselves, other Namespaces present in the Registry, and/or a replicable relative path structure and one or more file references from the targeted location.

The ability to Query the Namespaces Registry is required in all Environments where the Namespace Registry is supported. The Namespaces defining the infrastructure and Data Model versions supported in the Environment can be read, but entries can only be created and deleted by Administrators. There are no Namespaces Registry Events.

4.2 XML Schema Snippet

```
<xs:import namespace="http://www.w3.org/XML/1998/namespace"
schemaLocation="imports/xml/xml.xsd"/>
```

4.3 Supported Operations

Request	Direct Environment	Brokered Environment
Query	M	M
Create	P	P
Update	P	P
Delete	P	P
Events	P	P

There is no support for dynamic Query, and there are no defined Namespaces Registry Service Paths or XQuery Templates.

4.4 Namespaces Registry Data Elements

Element	Char	Description	Type
namespaces		Collection element used where multiple namespaces MAY be conveyed.	Collection
namespace	M		Namespaces Registry Entry
@id	M	Objects unique ID.	UUID
namespace/zone	M	If this namespace is tied to a specific zone it is specified here, otherwise "environment-global" indicates applicability to all zones.	xs:token
namespace/uri	MU	The namespace URI of the retrieved URL.	xs:anyURI minLength=0 maxLength=2048
namespace/url	MN	The URL of the specified namespace URI. If empty the namespace indicated by the URI is not valid within this Zone.	xs:anyURI minLength=0 maxLength=2048

4.5 XML Examples

```
<namespaces xmlns="http://sifassociation.org/infrastructure/3.0">
  <namespace id="97eea29f-013c-1000-007f-14109fdcaf83">
    <zone>environment-global</zone>
    <uri>http://sifassociation.org/infrastructure/3.0</uri>
    <url>http://sifassociation.org/infrastructure/3.0/core.xsd</url>
  </namespace>
  <namespace id="97eea29f-013c-1000-007f-14109fdcaf83">
    <zone>environment-global</zone>
    <uri>http://www.sifassociation.org/datamodel/us/3.0</uri>
    <url>http://www.sifassociation.org/datamodel/us/3.0/US3p00p04.xsd</url>
  </namespace>
  <namespace id="97eea29b-013c-1000-007f-14109fdcaf83">
    <zone>urn:edu:wwu</zone>
    <uri>urn:org:pesc:sector:AdmissionsRecord:v1.1.0</uri>
    <url>http://www.pesc.org/library/docs/standards/Admissions%20Application/AdmissionsRecord_v1.1.0.xsd</url>
  </namespace>
</namespaces>
```

5. Code Sets Registry

The SIF standard includes multiple normative dependencies on external code sets such as the NCES Handbook. When one of these code sets is revised, the first application utilizing the updated codes will likely not interoperate with previously deployed SIF applications conforming to the earlier version of the code set.

Any such a code set revision then represents a potential breaking change which is asynchronous to the SIF standard release cycle. As a result, Consumer / Provider interoperability is impacted even between SIF applications conforming to the same SIF minor release.

Even in the case where all Code Set values are defined directly in the SIF standard, the addition of a single new value to an enumerated code set in a minor release can break interoperability in ways that the addition of a new element cannot. Reception of a new element can be simply ignored by an application conforming to an earlier release, because everything it expected is still present in the arriving message. But reception of a new code set value cannot be ignored, because it provides a value which cannot be determined to be legal for an existing element that the older application may need to store.

The Code Set Registry Service provides a way for all legal codes to be defined outside of the SIF Specification while allowing changes (additions and replacements) of external code set values to be easily verified by the recipient so as not to break existing Consumer / Provider interoperability.

Codes are arranged (and scoped) under their respective code sets within the Registry. They are specifically not namespace qualified so they can be built up from multiple sources and remain simple to employ.

5.1 Service Implementation Strategy

Consumers and Providers can use the Code Set Registry to verify and then accept codes that were not yet defined when they were written. Depending upon the implementation, all or part of the Code Sets Registry can be seeded either manually or by having the Registry automatically draw in codes from multiple sources and reconcile them with the manual entries.

Query is a mandatory Code Set Registry operation. All others (*create*, *update* and *delete* of both code lists and their individual entries) are optional, and may require manual administrative actions to achieve.

However entered, if an unexpected code value is received by a Consumer or Provider, the Code Set Registry **must** be capable of verifying whether the value is in fact legal.

5.2 Alternatives

If a Code Sets Registry Utility Service is not available in the Environment, Consumers and Providers **MUST** employ one or more of the following strategies for dealing with the arrival of a code which does not match the expected enumerated list of values defined when they were created.

- **Blind Trust:** What is received is accepted by the recipient and placed into its data store. This is the simplest approach and removes the need to support or reference the Code Sets Utility Service. However it may result in inconsistent or erroneous data, since it ignores the possibility of sender error.

- **SIF File Import:** The SIF Association intends to publish a series of Code Set files in the format shown below. Administrators at a given site can deploy a Code Sets Registry and feed in these files as a way of initializing the Service.
- **Other Import Techniques:** Additional code sets can be copied directly from their source on the web or hand entered.

If the code value is not found in the Registry, the remainder of the message can be accepted, but either way an Alert **SHOULD** be issued.

5.3 Supported Operations

Request	Direct Environment	Brokered Environment
Query (both Bulk and Paged)	M	M
Create	P	P
Update	P	P
Delete	P	P
Events	M	M

There is no support for dynamic Query, and there are no defined Code Sets Registry Service Paths or XQuery Templates. Individual Code Set entries may be returned if Query by ID is requested.

5.4 Code Sets Registry Data Elements

Element	Char	Description	Type
codeSets		Collection element used where multiple code sets MAY be conveyed.	Collection
codeSet	M		Member
@id	MIU	The unique name of the code set	xs:token minLength=0 maxLength=128
codeSet/zone	M	If this is tied to a specific zone it is specified	xs:token

		here, otherwise “ <i>environment-global</i> ” indicates applicability to all zones in the Environment.	
codeSet/version	M	The official major, minor, and revision version of the code set. MUST be allowed to rev independent of (and more often then) both the data model and infrastructure used.	version
codeSet/timestamp	M	The date and time of the last change to this code set. So that only updates can be queried efficiently (where applicable).	xs:dateTime
codeSet/source	CN	URL to an external code set whose values are not (yet) present in the system. Either source or codeItems MUST be present, however both MUST NOT be included.	xs:anyURI minLength=0 maxLength=2048
codeSet/codeItems	C	Either source or codeItems MUST be present, however both MUST NOT be included.	
codeSet/codeItems/codeItem	MR	A code in the set.	
codeSet/codeItems/code	MI	Official abbreviation for the given value.	xs:token minLength=1 maxLength=16
codeSet/codeItems/source	SN	URL where the related abbreviation comes from.	xs:anyURI minLength=0 maxLength=2048
codeSet/codeItems/namespace	SN	Namespace where the related abbreviation comes from.	xs:anyURI minLength=0 maxLength=2048
codeSet/codeItems/value	SN	What the code represent.	xs:token

		Example: Morning Kindergarten	minLength=0 maxLength=128
codeSet/codeItems/description	ON	Human readable explanation of what the code represents.	xs:normalizedString minLength=0 maxLength=1024
codeSet/codeItems/definition	ON	Human readable explanation provided by the source!	xs:normalizedString minLength=0 maxLength=4096
codeSet/codeItems/aliases	O		
codeSet/codeItems/aliases/alias	MR		
codeSet/codeItems/aliases/alias/code	M		
codeSet/codeItems/aliases/alias/code/old	M	Indicates the alias code has been deprecated or deleted, even if it no longer appears in the register. Example: false	xs:boolean
codeSet/codeItems/aliases/alias/code/official	M	Indicates the alias code is an official one; however the locale uses this one to represent that concept, at least in part. Example: true MAY be used when cleaning data to leave the Environment.	xs:boolean
codeSet/codeItems/aliases/alias/code/value	M	Alternate abbreviation that might need to be more specific or correct.	xs:token minLength=1 maxLength=16
codeSet/codeItems/aliases/alias/source	SN	URL where the related abbreviation comes from.	xs:anyURI minLength=0 maxLength=2048
codeSet/codeItems/aliases/alias	SN	Namespace where the related abbreviation	xs:anyURI

/namespace		comes from.	minLength=0 maxLength=2048
codeSet/codeItems/action	M	The last action to be taken for this code item. Deleted code items SHOULD remain in the code set!	xs:enumeration ADD CHANGE DEPRECATED DELETE
codeSet/codeItems/timestamp	M	The date and time of the last action taken on this code item. So that only updates can be queried efficiently (where applicable).	xs:dateTime

5.5 XML Example

In the example below the SIF 2.x Data Model was used, and the appropriate Data Model Namespace is (counter-intuitively) “infrastructure/2.x”.

```
<codeSets xmlns="http://sifassociation.org/infrastructure/3.0" >
  <codeSet id="GradeLevels">
    <zone>environment-global</zone>
    <version>3.1</version>
    <timestamp>2012-09-01T07:00:13.000-07:00</timestamp>
    <codeItems>
      <codeItem>
        <code>K</code>
        <source>http://vocabulary.curriculum.edu.au/schoolLevel</source>
        <namespace>http://www.sifinfo.org/infrastructure/2.x</namespace>
        <value>Kindergarten</value>
        <action>ADD</action>
        <timestamp>2012-01-01T07:00:13.000-07:00</timestamp>
      </codeItem>
      <codeItem>
        <code>1</code>
        <source>http://vocabulary.curriculum.edu.au/schoolLevel</source>
        <namespace>http://www.sifinfo.org/infrastructure/2.x</namespace>
        <value>First Grade</value>
        <action>ADD</action>
        <timestamp>2012-01-01T07:00:13.000-07:00</timestamp>
      </codeItem>
    </codeItems>
    <!-- Ten more entries here (2nd through 11th). -->
    <codeItems>
      <codeItem>
        <code>12</code>
        <source>http://vocabulary.curriculum.edu.au/schoolLevel</source>
        <namespace>http://www.sifinfo.org/infrastructure/2.x</namespace>
        <value>Senior</value>
        <action>ADD</action>
        <timestamp>2012-01-01T07:00:13.000-07:00</timestamp>
      </codeItem>
    </codeItems>
  </codeSet>
</codeSets>
```



```
</codeItem>
<codeItem>
  <code>KA</code>
  <value>Morning Kindergarten</value>
  <aliases>
    <alias>
      <code>
        <old>false</old>
        <official>true</official>
        <value>K</value>
      </code>
      <source>http://vocabulary.curriculum.edu.au/schoolLevel/</source>
      <namespace>http://www.sifinfo.org/infrastructure/2.x</namespace>
    </alias>
  </aliases>
  <action>ADD</action>
  <timestamp>2012-09-01T07:00:13.000-07:00</timestamp>
</codeItem>
</codeItems>
</codeSet>
</codeSets>
```

6. XQuery Templates Registry

The XQuery Templates Registry contains the Environment-wide collection of registered XQuery Templates (XQuery scripts which incorporate externally set parameter values). This collection defines the entire set of XQuery scripts that Consumers may legally issue for execution by Service Providers. Each XQuery Template has an associated unique ID under which it may be referenced.

Many if not all of the templates in a typical XQuery Templates Registry may have originally been specified as part of the Infrastructure “binding” of the Data Model release. For example an XQuery Template ID of “StudentSnapshot” will, when requested, result in a Response with a predefined format. Such templates are known “in advance” by developers of both Service Consumer and Service Provider applications, and serve primarily as a report documentation tool. Exactly how the report is generated remains an implementation detail left to the Service Provider.

Other XQuery Templates may be installed specifically by the site administrator or dynamically created by administrative-level Consumer applications.

All are used identically. A Consumer specifies the ID of an XQuery Template in a Query Request to a specific Service Provider, along with a set of values for any associated script parameters⁸.

The Provider may already either:

- Have its own copy of the script template corresponding to the supplied ID. In this case it proceeds to execute the script and return the results as the Response to the Query (either bulk or paged).
- Have its own logic ready to generate the expected results in the Response. When true, it allows XQuery Tokens in a Request to be serviced by a Provider which neither knows nor supports XQuery technology.

Otherwise the XQuery Template is unknown to the recipient Service Provider. In order to successfully process such a Consumer Query, the Provider must first issue its own Query by ID to this XQuery Templates Registry to acquire the XQuery Template corresponding to the supplied ID. The Service Provider **should** then:

- Cache this Template (so it can be reused for later Requests)
- Execute it as an XQuery script by applying the associated parameters
- Pack the results into the Response.

If the Provider is unable to support this logic, than a SIF Error Message **must** be returned with an HTTP Error Code of 405 (Method Not Allowed).

6.1 Security

<** Note: This entire section relates solely to Consumer-supplied XQuery Templates and can be ignored if all entries in the XQuery Template Registry are predefined by Site Administrators, guided in part by the collection of XQuery Templates contained in the binding of the SIF Data Model deployed at the site.>

⁸ Please refer to the XQuery section in the Base Architecture Document for further details.

The XQuery Templates Registry can optionally enforce administrator security over the formats of acceptable XQuery Scripts which may be utilized in the Environment. For example, any script in which element names are altered or element values combined and reported as the value of another element **SHOULD** be prohibited. This removes the possibility that the presence of XPath-based XML filters between Consumer and Provider in accordance with site security policies can be circumvented by the internals of an XQuery script.

6.1.1 *Lazy Authorization*

There may be a significant delay between the time a Consumer attempts to create a new XQuery Template in the Registry and the time when it is approved, since the XQuery Template contained in the *Create Request* may have to be reviewed by administrative personnel.

To determine whether or not the Template has been approved, the Consumer can issue a Query by ID for the Template it just created. This will contain a *status* element with one of three possible Responses:

- **Pending.** The Consumer should reissue the Query at timed intervals until one of the other statuses is returned.
- **Disallowed.** The Template was rejected. The *qualifier* element will contain the reason.
- **Approved.** The Template has been accepted, and can be referenced in subsequent XQueries to other Service Providers

If the status is “*Pending*” the Consumer can either re-query at a fixed interval until the status changes, or alternatively (since the Status change represents a modification to the Template), the Consumer can wait for the corresponding Event for that XQuery template, and determine whether it reports a change from “Pending” to “Allowed” or Disallowed”.

6.1.2 *Security Limitations*

XQuery Templates are guaranteed to be secure only in their script limitations, and **NOT** in the particular instance in which they may be executed. There may be external administrator-imposed limitations on which Consumers can utilize a specific template, and limitations on the Zones in which the recipient Provider can reside.

For those dynamic aspects, the same authorization limitations that apply to a specific Consumer issuing a Paged Query to a specified Service Provider (which attempts to retrieve every element of every object) apply to any XQuery script the Consumer may issue to that Provider as well.

As a result, there is no information associated with an XQuery Templates Registry entry that relates to a specific Consumer, Zone or Context.

An authorized Consumer **MAY** attempt to create XQuery Templates for later use by itself or other Consumers. Once approved, the XQuery may be utilized in Queries by any authorized Consumer with the Template ID.

Upon receiving a Create Request the XQuery Template Registry Provider sets the following items:

- **id:** **MUST** use the provided value. If it is a duplicate of an existing XQuery Template, the *create Request* **MUST** be rejected.

- **type:** If the contained XQuery Template is limited to requesting all or a subset of defined object elements in all or a subset of qualifying objects of a single type, the Template type **MUST** be set to SINGULAR.
- If the XQuery Template does not fit the Singular format, but is still limited to a single object type (as when it also includes requesting “calculated” values based upon existing elements (ex: “the sum of the number of students who ...”) the Template type **MUST** be set to FORMULA.
- Everything else (XQuery Templates containing references to objects of more than one type) **MUST** be set to EXTENDED.
- **status:** Object Providers **MUST** only execute XQueries with a status of APPROVED. This status **MAY** be assigned in a variety of ways (blind trust, manual administrator approval, or automated XQuery Template analysis).

As the Environments Provider encounters messages that reference XQuery Templates, it **MAY** enforce restrictions through ACLs. For example: allow constructs such as:

```
reference /sd:xquery[id="StudentsBySchoolName"]
```

When a Consumer invokes a Query which includes the ID of a registered XQuery Template, the Service Provider receiving it **MUST** accomplish the following.

- Resolve the Template ID to the contents of its Registry entry. For performance reasons, Providers **SHOULD** stay synchronized with the XQuery Templates Registry, but **MAY** Query for the Template to resolve each arriving Query.
- Ensure the entry is of a *type* supported by this provider and its *status* is APPROVED.
- Ensure all parameter values contained in the Request have corresponding entries in the XQuery template and are valid for use.
- Substitute parameter values into the XQuery Template to make the XQuery Script
- Execute the script and verify that accurate results have been secured.

If an include statement is part of the registry entry, the provider **SHOULD** reduce elements of the resulting XML down to those specified, before sending results to the consumer. Note that especially when working with small objects; it is generally more efficient to receive the entire data object than to ask the Provider to return a partial object.

A Consumer can dynamically construct an “ad-hoc” XQuery in response to some user-selected form by creating, using, then deleting the need XQuery Template from this registry.

6.2 Restrictions on Use

Singular Form

This more limited single object Query operation only supports queries that align with the Singular XQuery Template capabilities defined below. Note that it is a superset of the dynamic Query “*where*” URL Query Parameter that is defined in the Basic Architecture.

- Each element in each data model reference must be restricted to the actual (qualified) name of the entity being referenced. In other words, no wildcards (`//`, `..`, `..*`, `@*`, `node()`) may be used in the XPaths.
- XQuery and XPath functions **must not** be used in a Singular XQuery.
- Variables **must not** be used in a Singular XQuery.
- Singular XQueries **must** match the structure shown here (excluding whitespace).
- Transports may convey the Singular XQuery in the format most suitable for them; however that format **must** map directly to the structures defined in the following table.
 - Namespace_Declarations
 - Entity_Name[

(Element Operator Value Inner_Boolean_Operator Element Operator Value) Outer_Boolean_Operator

(...)]

Component	Examples	Comments
Namespace Declarations	declare namespace dm = "http://www.sifassociation.org/datamodel/na/3.1 declare default element namespace " http://namespaces.sifassociation.org/compliance/csq ";	As individual Data Entities may contain multiple namespaces, multiple Namespace Declarations must be supported.
Entity Name	/dm:StudentPersonal /CombinedCSQ	This both indicates what Object to return as well as the root element of the XPaths located in the where clause (if any).
Where Clause	[...]	
Outer Boolean Operator	“and,” “or,” or omitted	May use as many as possible, they all must be the same.
Condition Group	(...)	
Inner Boolean Operator	“and,” “or,” or omitted	May use as many as possible, they all must be the same.

Condition	dm:Name/dm:LastName="Smith" ObjectMatrix/Object/@processesAdds = "true"	Logical expression bound by an operator.
Element	dm:Name/dm:LastName ObjectMatrix/Object/@processesAdds	XPath to the element (or attribute) being evaluated.
Operator	= < > <= >= !=	Transports may require these or other characters be escaped on the wire as is appropriate to the technologies employed.
Value	"Smith" "true"	XQuery requires non-numeric values to be surrounded with double quotes (and vice versa).

The XQuery Templates Registry can, by confirming the exclusive presence of the following lines (in order), verify that the Singular XQuery form is being followed:

- One or more XQuery namespace declarations.
- One restricted XPath 2.0 expression + parameters.

Example:

The following example is of a Singular XQuery script, which searches for all students with a particular surname.

```
declare namespace dm = "http://www.sifassociation.org/datamodel/us/3.0";

/dm:Student [(dm:Name/dm:LastName="Smith") and (...)]
```

If the XQuery returned no qualifying objects, the Service Provider should return an HTTP error code of 404 (Not Found). Otherwise the XQuery returns one or more Student objects in the form.

```
<student> ... </student>
<student> ... </student>
```

Before putting this result out on the wire as the Response, a Service Provider **must** “wrap” these results in a <students> element to conform to the formats returned by the other types of Query (Data Object and Service Path).

```
<students>
  <student> ... </student>
  <student> ... </student>
</students>
```

Formula Form

SIF intends a Formula to be an XQuery that does more than just retrieve data; it may do calculations and other useful transformations to greatly increase efficiency (over the wire) as long as namespace and other security rules are observed. Formally a Formula is any XQuery that does not follow the singular form, however still only targets a single data object. That said application specific function call are discouraged and any XQuery containing them **should** be rejected by the Providers registry.

Extended Form

Extended Queries are similar to Formulas however they may target and/or combine multiple objects.

Parameters

While named parameters provide the ability to pre-approve Templates where one or more values differ between Requests utilizing the same Template, Providers must be able to ensure these values specified are not more than a single value. To that end only alphanumeric values (including floating point numbers) with spaces or the empty string are allowed as parameter values. This restriction **MAY** require what is conceptually one value to be communicated through multiple parameters. For instance: “{.ns:}{.object:}”

Access and Ownership

Depending upon XQuery Template Service functionality and a Consumer’s authorization rights, a Consumer may create a Template, query all Templates in the Registry by either Paged Query or Query by ID (the template token) and delete only those Templates it has created.

6.3 Supported Operations

In most cases the contents of the XQuery Template Registry will be fixed, and only the Query operation will be supported.

Request	Direct Environment	Brokered Environment
Query	M	M
Create	O	O
Update	P	P

Delete	O	O
Events	O	O

There is no support for dynamic Query, and there are no defined XQuery Templates Registry Service Paths or XQuery Templates which self-reference this Service. Individual XQuery Templates may be returned if Query by ID is requested.

6.4 XQuery Templates Registry Data Elements

Element	Char	Description	Type
xquerys		Collection element used where multiple XQuerys MAY be conveyed.	Collection
xquery	M		Member
@id	M	The unique ID of the XQuery Template, which is contained in the Query Request URL when this template is being utilized. “Suggested” by the Consumer issuing the <i>create</i> Request, or pre-set in accordance with the binding requirements of a particular SIF release. This is NOT required to be a UUID, and must be provided as the Service Identifier in every Query URL which invokes this Template. Ex: “StudentSnapshot”.	xs:token
/type	M	Determines if the query may be sent to a data provider or if it requires the extended XQuery mechanism. Determined by the XQuery Templates Registry Service	<u>enumeration</u> SINGULAR FORMULA EXTENDED
/status	M	Communicates if the query has been inspected for data hiding and the results of that assessment. Determined by the XQuery Templates Registry Service Only APPROVED XQueries MAY be executed.	<u>enumeration</u> PENDING APPROVED DISALLOWED
/qualifier	OC	If status is DISALLOWED this element optionally contains the reason. Determined by the XQuery Templates Registry Service	xs:normalizedString minLength=0

			maxLength=1024
/description	ON	Human readable description of the XQuery's intended behavior. Example: Query all the students in SIF University.	xs:normalizedString minLength=0 maxLength=1024
/script	MN	The text of the XQuery which MAY accept parameters. By replacing fields surrounded by { : and : } with the value with the name of the field. XQueries are capable of hiding data. Therefore the Template SHOULD be scrutinized against privacy policies before being permitted to run.	CDATA
/returnType	Q	The fully qualified name of the root element returned by this service. Example: http://www.sifassociation.org/datamodel/us/3.0:students	xs:anyURI

Example

In the example below the SIF 2.x Data Model was used, and the appropriate Data Model Namespace is (counter-intuitively) "infrastructure/2.x".

```
<xquerys xmlns="http://sifassociation.org/infrastructure/3.0">
  <xquery id="StudentsBySchool">
    <description>Query all the students in the specified school.</description>
    <script>
<![CDATA[
(: Query all the students in the school (i.e. name=SIF University). :)
declare namespace dm = "http://www.sifinfo.org/infrastructure/2.x";
for $school in /dm:SchoolInfo[dm:SchoolName = "{:name:}"]
for $enrollment in /dm:StudentSchoolEnrollment
for $student in /dm:StudentPersonal
where $school/@id = $enrollment/@SchoolInfoId and $enrollment/@StudentPersonalId =
$student/@id
return $student
]]>
    </script>
  </xquery>
</xquerys>
```

7. Alerts

There is a single Alerts Utility Service available in the Environment, located in the *environment-global* Zone. It replaces and extends the functionality provided by the LogEntry Event in SIF v2.x, providing authorized Service Consumers with the ability to:

- Log an “alert” (error, warning or status change) by creating an Alert object
- Retrieve specified collections of these Alert objects via standard query mechanisms
- Subscribe to Alert creation Events

The ability to detect when new Alerts are created could be used by an administrative Consumer to monitor the performance of newly installed applications, detect when another Consumer had indicated it was the cause of a problem, or track and flag Alerts above a pre-specified priority level. Such Alerts might include preventative maintenance warnings from a Queue Infrastructure Service that its Consumer has stopped polling for arriving messages, or that the number of messages in the Queue have crossed a predefined threshold.

Ideally an Alert should contain as much identifying information about the problem being reported as possible. However “nesting” an erroneous message inside the Alert can generate unanticipated problems if the error being reported is that the original message format was invalid. For this reason, the original message **MAY** be omitted or described and when it is included it **MUST** be properly escaped (included as CDATA).

7.1 Supported Operations

The required levels of supported Alerts Utility Service operations (both Requests and Events) available to properly authorized Consumers are indicated in the table below (where M = Mandatory, O = Optional, P = Prohibited).

Operation	Direct Environment	Brokered Environment
Query	O	O
Create (single object form only) ⁹	M	M
Update	P	P
Delete	P	P
Event	M	M

⁹ Only one Alert object can be specified in each Alert Create Request.

Support for Query is optional, and when provided for non-Administrative-level applications, returns only Alert objects created the same application that issued the Query. There is no support for dynamic Query, and there are no defined Alert Service Paths or XQuery Templates. Individual Alerts may be returned if Query by ID is requested.

7.2 Error Handling

Proper use of the Alerts Service is an essential part of the error logic handling of every Service Consumer and Provider. The following message exchange situations during which errors occur define how SIF 3.0 components **should** interact with the Alert Utility Service, and are applicable to all Service types (Object, Functional and Utility).

7.2.1 A Service Provider receives a Consumer operation invocation which it rejects

Actors	Consumer which invokes a Service operation Functional or Data Entity Service Provider supporting that operation Alerts Utility Service
Preconditions	The Consumer invokes a Service operation. It reaches the Service Provider, which rejects it. The reasons could range from: <ul style="list-style-type: none"> • Invalid XML in the Request payload • Omission of a mandatory element • Incorrect data value (ex: an Object ID which did not correlate to a stored object)
Actions	The Service Provider should then perform the following actions: <ul style="list-style-type: none"> • Return a “NAK” Response to the client, indicating at least the error category, code and description which identify the problem. • Optionally determine whether this error is a duplicate or deserves to generate a unique Alert (i.e. is it identical to a previous error about the same Consumer that has already been logged?). If not, ignore it (or file a different Alert). • If applicable, the reporting Service should issue a <i>Create Alert</i> Request to the Alerts Utility Service, identifying itself as the Reporter and the Consumer as the Cause, and providing as much information about the error as seems reasonable. <p>This should indicate at least the description, error, category and code which document the problem in the Consumer Request. Along with the appropriate error category and code:</p> <ul style="list-style-type: none"> ○ For invalid XML, return the error reported from the parser ○ For omission of a mandatory element, return the element tag name ○ For an invalid data value, return the element tag name and the erroneous value • Take any other action (ex: logging the error to a local file) as seems appropriate
Post Conditions	When receiving the rejection of its Service method invocation, the Consumer should then perform only those actions which relate to its own internal logic (i.e. log the problem to a local file, report it

	<p>to the user, back out of a transaction, adjust its internal data base).</p> <p>However the Consumer does not normally create an Alert object reporting the problem, as it can rely upon the Service Provider to do that.</p>
--	---

7.2.2 *The Consumer receives a Service Provider Response which it rejects*

Actors	<p>Consumer which invokes a Service Provider operation</p> <p>Functional or Object Service Provider supporting that operation</p> <p>Alerts Utility Service</p>
Preconditions	<p>The Service Provider Response to a previously invoked operation is seen and rejected by the Consumer. The reasons could range from:</p> <ul style="list-style-type: none"> • Invalid XML in the Response payload • Omission of a mandatory element • Incorrect data value (ex: an enumerated value such as a Country Code which does not correspond to a valid entry as the Consumer understands it, in an External Code List)
Actions	<p>The Consumer has no way to report this problem back to the Service Provider. It should then:</p> <ul style="list-style-type: none"> • Perform only those actions which relate to its own internal logic (i.e. log the problem to a local file, report it to the user, back out of a transaction, adjust its internal data base). This closely parallels the required actions in the case when the Service Provider has rejected the Consumer's Request. • If the Consumer believes the cause of the disconnect rests with the Service Provider, it should also issue a <i>Create Alert</i> Request to the Alerts Service, identifying itself as the Reporter and the Service Provider which returned the Response as the Cause, providing as much information about the problem as seems reasonable. <p>This should indicate at least the description, error, category and code which documents the problem in the Response. Along with the appropriate error category and code:</p> <ul style="list-style-type: none"> ○ For invalid XML, return the error reported from the parser ○ For omission of a mandatory element, return the element tag name ○ For an invalid data value, return the element tag name and the erroneous value
Post Conditions	<p>In rare cases, the Consumer may "re-request" the operation using alternative parameters if it has reason to anticipate that will produce better results.</p>

7.2.3 *A Service Provider posts an Event which a Subscribing Consumer rejects*

Actors	<p>Consumer Subscriber to Service Provider Events</p> <p>Service Provider which publishes a Change Event</p>
---------------	--

	Alerts Utility Service
Preconditions	<p>An Event is received and rejected by the Subscriber. The reasons could range from:</p> <ul style="list-style-type: none"> • Invalid XML in the Event payload • Omission of a mandatory element (in a Create Event only) • Incorrect data value (ex: an enumerated value such as a Country Code which does not correspond to a valid entry as the client understands it, in an External Code List)
Actions	<p>The Subscriber has no way to report the problem back to the Service. It must ignore the Event in terms of processing it or updating its Data Store. However it should log the problem to a local file and / or report it to its end user.</p> <p>There is also the possibility that an erroneous Provider is publishing a stream of faulty Events, which if left unchecked, would result in a flood of new Alert Objects from each Subscriber. As a result it is recommended that before a Subscriber creates an Alert to report an Event error, it first determines that it has never previously (or not in a pre-specified time) reported an error contained in an Event of this type.</p> <p>If the Subscriber believes the cause of the disconnect is a problem with the Publisher, and if it determines it is not redundant to do so, it should also issue a <i>Create Alert</i> Request to the Alert Service, identifying itself as the Reporter and the Service which published the Event as the Cause, providing as much information about the original Event as seems reasonable.</p> <p>This should indicate at least the description, error, category and code which documents the problem. Along with the appropriate error category and code:</p> <ul style="list-style-type: none"> • For invalid XML, return the error reported from the parser • For omission of a mandatory element in a Create Event, return the element tag name • For an invalid data value, return the element tag name and the erroneous value
Post Conditions	<p>The Subscriber should record that it has created an Alert because of an erroneous Event received for this Object type. This will help minimize the number of Alerts reporting the same problem should the cause reside with the Event Publisher.</p>

7.2.4 A Service Provider detects an inactive Consumer

Actors	<p>Stateful Service Provider dependent on additional Consumer operations</p> <p>Alerts Utility Service</p>
Preconditions	<p>The Service Provider detects a Consumer error because an expected Service operation has not been invoked.</p> <p>An example would be when a Polling Queue Instance determines that the preset maximum time limit to wait for the arrival of a GetNextMessage invocation has expired, indicating that the Consumer may be offline.</p>
Actions	<p>The Service has no opportunity to report the problem back to the Consumer (and in fact the</p>

	<p>Consumer may not even be active).</p> <p>The Provider should then perform the following actions:</p> <ul style="list-style-type: none"> • Determine whether this error is a duplicate or deserves to generate a unique Alert (i.e. is it identical to a previous error about the same Consumer that has already been logged?). In the example given, if the error is deemed duplicative, it should be ignored. • If applicable, the reporting Service should issue a <i>Create Alert</i> Request to the Alerts Utility Service, identifying itself as the Reporter and the Consumer as the Cause, and providing as much information about the error as seems reasonable. This should indicate at least the description, error, category and code which document the “state” problem the Service is having with the Consumer <p>When reporting “<i>Consumer inactivity</i>”, along with the appropriate error category and code, the Alert object created should contain Alert Level, Error and Description element values which provide indications of:</p> <ul style="list-style-type: none"> ○ Time of last client activity ○ Size (where known) of current Message Queue ○ Severity Level of Alert (warning that limit exceeded or application offline error) <ul style="list-style-type: none"> • Take any other action (ex: logging the error to a local file) as seems appropriate
--	---

7.3 Alert Object Data Elements

The following data elements comprise the payload of an Alert Object “create” request. Note that there is no “Id” attribute as that is supplied by the Alerts Utility Service Provider when the Alert Object is actually created.

Element or @attribute	Char	Description	Type
alert		Top level element of an Alert	Alert Object Root
alert@id	M	Identifier for the <i>alert</i> object. Defined by the Alerts Service when a create Alert is issued.	UUID
alert/reporter	M	External identification of the Application (Consumer or Provider) reporting this Alert. This is typically the Consumer Name of the reporting application.	xs:token
alert/cause	O	External identification of the cause of the Alert. This is typically the <i>sourceName</i> of the Partner as contained in the HTTP Header Field of the message which provoked the creation of the <i>alert</i> .	xs:token
alert/exchange	M	The exchange (or lack of exchange) responsible for generating	xs:enumeration

		the Alert.	One of: REQUEST RESPONSE EVENT TIMEOUT
alert/level	M	The level of the Alert.	xs:enumeration One of: INFO STATECHANGE WARNING ERROR
alert/description	O	A description of the reason for the Alert	xs:normalizedString
alert/messageId	C	If available, the ID of the Message causing the problem	xs:token
alert/body	O	The internals of the offending message or a more complete description of the information or state change.	CDATA
alert/error	O	Detailed error results such as a stack trace	xs:string
alert/xpath	O	An indicator of the specific element that was in error (or contributed to the problem).	xs:selector
alert/category	MC	If the Alert Level is an error, this value must be the SIF 3.0 Error Category ¹⁰ corresponding to the type of Error being reported. Otherwise it should be mapped to a corresponding category type where possible. If this message was persisted and processed later, the corresponding Category MAY be included here, even though the message was originally acknowledged normally.	xs:unsignedInt
alert/code	MC	If the Alert Level is an error, this value must be the SIF 3.0 Error Code ¹¹ corresponding to the type of Error being reported. Otherwise it should be mapped to a corresponding category type	xs:unsignedInt

¹⁰ Please refer to Appendix D of the Infrastructure Services document.

¹¹ Please refer to Appendix D of the Infrastructure Services document.

		<p>where possible.</p> <p>If this message was persisted and processed later, the corresponding Code MAY be included here, even though the message was originally acknowledged normally.</p>	
alert/internal	O	Code internal to the reporter.	xs:token

7.4 XML Examples

The payload of an Alert issued by a Data Miner regarding data returned from and SIS system.

```

<alert id="">
  <reporter>49erDataMiner</reporter>
  <cause>RamseySIS</cause>
  <exchange>REQUEST</exchange>
  <level>ERROR</level>
  <description>Date format not understood.</description>
  <body>
    <![CDATA[
      <StudentSchoolEnrollment      Id="A8C3D3E34B359D75101D00AA001A1652"
      StudentPersonalId="D3E34B359D75101A8C3D00AA001A1652"
      SchoolInfold="D3E34B359D75101A8C3D00AA001A1651"
      MembershipType="Home" TimeFrame="Current" SchoolYear="2004">

      <EntryDate>2004-01-29</EntryDate>
      <EntryType>
        <Code>1838</Code>
      </EntryType>
      <GradeLevel>
        <Code>10</Code>
      </GradeLevel>
      <Homeroom
        SIF_RefObject="RoomInfo">D7510D3E34B3591A8C3D00AA001A1651
      </Homeroom>
      <Advisor
        SIF_RefObject="StaffPersonal">B359D3E34D75101A8C3D00AA001A1652</Advisor>
      <FTE>1.00</FTE>
      <FTPTStatus>FullTime</FTPTStatus>
      <ResidencyStatus>
        <Code>1653</Code>
      </ResidencyStatus>
      <NonResidentAttendReason>1658</NonResidentAttendReason>
      </StudentSchoolEnrollment>
    ]]>
  </body>

  <error>Missing tag "ExitDate."</error>
  <xpath>

```



```
<path>/StudentSchoolEnrollment/ExitDate</path>
<namespaces>
  <namespace>
    <prefix nil="true"/> <!-- nil = default -->
    <static>http://www.sifinfo.org/dataModel/3.0</static>
  </namespace>
</namespaces>
</xpath>
<category>1</category>
<code>4</code>
</alert>
```

Appendix A: Minimal (Simple) Consumer

The minimum (or “Simple”) SIF 3.0 Consumer is typically a REST-based application running on a mobile device that needs to initiate data exchanges with one or more data providers.

The decision to code such an application as a SIF Consumer ensures its developers that it will interoperate securely and robustly “out of the box” with a broad set of products that supply SIF-compliant object data. This appendix details the extent of “additional interactions” such minimum SIF Consumers **must** have with their provided Environment, above and beyond that of a RESTful client issuing Query, Create, Update and Delete requests to one or more RESTful Services.

The following steps reflect the complete set of mandatory requirements imposed on even the simplest SIF 3.0 conformant client application. It’s as easy as 1-2-3.

1. Register as a SIF 3.0 Consumer by issuing a *create* Request (HTTP or HTTPS POST) to a pre-set *Environments* Service URL, using either the BASIC or HMACSHA256 authentication method.
2. Issue one or more SIF 3.0 compliant single object requests (to the Request Connector Infrastructure Service URL that was returned in the payload of the Environment *create* Response), and process all object responses.
3. Issue a “*create*” request to the Alerts Utility Service if / when any unexpected errors occur.

And that’s it! The minimal Consumer represents the first step on the SIF 3.0 adoption ladder. From this simple beginning, a wide range of additional functionality can be seamlessly added. Some examples are shown below.

Intermediate Consumers:

- Subscribe to and receive asynchronous Events published by selected Service Providers
- Pack multiple objects in each request
- Qualify Query Requests by adding a dynamic “where” and clause, and specify the “page” of results which is desired.
- Qualify Query Requests by using a Service Path (ex: ../schools/1234/students to obtain all students in school 1234)

Advanced Consumers:

- Examine the full range of Service Providers accessible through their Environment, and self-provision themselves to a set of additional Services, possibly spanning multiple Zones and Contexts
- Construct and use “formula” XQuery scripts to calculate and return additional elements
- Register as a Provider of one or more Services accessible by other Consumers (in Brokered Environments only).

The remainder of this Appendix further defines the Minimal Consumer, and details the full set of Infrastructure and Utility Service functionality it uses.

A.1 Minimal Consumer vs. more “demanding” Consumers

The following table contrasts aspects of the Minimum Consumer with the more fully functional Consumer typical of SIF v2.x deployments.

Aspect	Minimal SIF 3.0 Consumer	SIF 3.0 Consumer equivalent to those that functioned as SIF 2.6 Clients
Example	“Portal-type” Consumer application which allows individual students to access their personal day’s schedule and class assignments.	Student Portal Application which can access everything related to a student and uses student identity to determine what will be reported
Typical Deployment	Mobile Device	Data Center
Required Environment Functionality	Consumer application requires functionality equivalent to a simple browser-based GUI (no Events, single object requests only, no XQuery, assumes successfully provisioned). Application instance can register and successfully function as a Consumer in any SIF 3.0 Environment, “Minimal” or above.	Consumer application can leverage extensive SIF 3.0 functionality (Events, multiple objects per request, XQuery and paged queries, ability to discover available services, ability to self-provision, ...) Application cannot register as a Consumer in a Minimal Environment, because it exploits functionality beyond the bare minimum supported there
Numbers	Many simultaneously registered instances, each representing one student	One or a few registered instances, covering all students.
Security	Authorization rights are based upon who is the User of the Consumer Application, and their role within the educational organization. The User ID, based upon User Name and Password, may form the basis of the security scheme.	Authorization rights are pre-assigned to the Consumer Application. The Application ID, in conformance with the described Consumer registration authentication elements, will form the basis of the security scheme.
Duration of a typical “Session”	Short. When the student signs out, the Environment is deleted	Long. The Environment, once created, is generally maintained indefinitely.

A.2 Core Infrastructure Service Interactions

The required interactions between a Simple Consumer and each of the set of SIF 3.0 core Infrastructure Services are contained **in bold** in the table below. They amount to one additional Create and a read of the information contained in the Response.

All other interactions (including anything past the 3rd service below) are optional.¹²

Basic Infrastructure Services	Minimal Consumer Usage	Details and common Consumer enhancements
Environments	<p>The Consumer must issue a create request to the <i>environments</i> URL containing the application's identification and its transport requirements.</p> <p>On success, it gets back its "<i>environment</i>" and a Session Token which must be used on all subsequent Service Requests.</p>	<p>This single operation invocation registers the application as a SIF Consumer, and assigns it an Environment from which all needed Services can be accessed.</p> <p>The Environment could be one of several previously defined for the site (ex: Testing, Staging, Production), depending on the Application registering as a Consumer.</p>
Environment	<p>The <i>environment</i> returned in the <i>create</i> response has the URLs of the other Infrastructure Services comprising and supporting it, along with the Consumers default Zone ID and initial service operation authorization rights.</p> <p>The Consumer must save the URL of the Requests Connector Infrastructure Service</p>	<p>The data contained in the Consumer's Environment does not change once Registration is complete, except by Consumer-specific action (like creation of a <i>provisionRequest</i>).</p> <p>As a result, once queried, it does not have to be queried again.</p>
Requests Connector	All Data Object and Utility Service operations are invoked here.	This service is treated identically to a RESTful Service URL. There are no Connector-specific operations for the Consumer to request.
Queues	<p>A minimal Consumer will issue only "immediate" Requests, which result in synchronous Responses being returned on the same HTTPS connection (nothing is queued). Only Events will arrive asynchronously.</p> <p>If Object Events are not subscribed to, there is no interaction with the Queues Service.</p>	<p>If a Consumer does intend to subscribe to object Events, it must first issue a "Create Queue" request to the Queues Service.</p> <p>This Queue will be used to receive asynchronously arriving Event messages.</p>
Provision Requests	If site administration policies support "pre-authorization" of access rights (as indicated in the returned Environment), the Consumer can assume its environment already contains pre-authorization of all the Requests it needs.	<p>If the Environment assigned to the application does not support pre-authorization", then the Consumer must either:</p> <ul style="list-style-type: none"> • Issue a "Create" Request to the <i>ProvisionRequest</i> Service to attempt to

¹² All the Infrastructure Services contained in this table are defined in far more detail in the Infrastructure Services document.

		<p>manually acquire the set of authorization rights it needs to access its selected services.</p> <ul style="list-style-type: none"> Just issue the request and respond to a “non-authorized” Error Response by creating an <i>Alert</i> (the suggested minimum Consumer strategy).
Subscriptions	If the Consumer does not subscribe to Events, it has no interactions with this service.	<p>After a (non-minimal) authorized Consumer has created a Queue to receive published events, it needs to <i>create</i> a Subscription to associate that Queue with the Service Provider publishing them.</p> <p>Then the Consumer must issue a timed series of synchronous “GET” (poll) requests to the Queue to actually obtain those published Events.</p>

A.3 Utility Service Interactions

The required interactions between a Simple Consumer and each of the set of SIF 3.0 Utility Infrastructure Services are contained **in bold** in the table below.¹³ They amount to using the Alerts service to effectively throw exceptions if unexpected errors are encountered.

All other interactions (including anything past the 1st service below) are optional

Utility Service	Simple Consumer Usage	Details and common Consumer enhancements
Alerts	Any Consumer encountering an unexpected or erroneous condition must issue a <i>create</i> Request to the Alerts Utility Service to create an Alert reporting the problem it encountered.	Every Consumer must be authorized to issue <i>create</i> Alert Requests to the Alerts Utility Service.
Zones Registry	If a Consumer only interacts with the set of Service Providers contained in its own assigned “default” Zone or the available Utility Services in the <i>environment-global</i> Zone, it does not need to interact with the Zones Registry.	The Consumer may use this Registry to optionally explore the set of Zones (and associated Service Providers) made available to it through its supplied Environment.
Providers Registry	If a Consumer already knows the name of the	The Consumer can optionally verify the

¹³ All the Utility Services contained in this table are defined in far more detail in the earlier sections of this Utility Services document.

	object types (ex: Student) and optionally, the Contexts (ex: "Current") it will access, and if it is going to assume they are located in its default Zone, it does not need to interact with the Providers Registry.	presence of any targeted Object Provider (or fail to do so) and if successful, confirm additional details about it (including the Context of the objects it provides).
Namespaces Registry	Any Consumer can adopt the "blind trust" strategy and accept the XML payloads it receives without validating it against pre-stored schemas (i.e. run with validation off). In this case it does not need to interact with the Namespaces Registry.	This Utility service allows the validating recipient of an XML message to avoid rejecting the payload of a message containing elements defined in a properly registered SIF Data Model extension.
Code Sets Registry	Any Consumer can adopt the "blind trust" strategy, and assume any Code Set value it does not understand is legal, and accept it. If it does this, it does not need to interact with the External Code Sets Registry.	The Consumer (or Provider) can optionally attempt to verify whether an "unexpected" code set value it has received is simply unknown to it or illegal, by querying the External Code Set Registry.
XQuery Templates Registry	Any Consumer which restricts itself to issuing either unqualified queries or queries qualified only by URL parameters, does not need to interact with the XQuery Template Registry. A slightly more featureful Consumer might issue "static" XQuery Requests in which the XQuery Template name was pre-defined by the Data Model it supported. In this case, the XQuery Template would already be pre-registered, and Consumer would again not have to access the XQuery Template Registry.	Only those Consumers who dynamically create XQuery Templates need access the XQuery Template Registry.

Appendix B: Minimal Environments Providers

The “Minimal Environments Provider” supports only the set of features that every client application can rely on being present, once it has successfully registered itself as a Consumer in a SIF 3.0 Environment. Additional functionality can be (and often will be) optionally supported by a given Environment Provider, but the minimal requirements described below are mandatory for all of them.

B.1 Overarching Design Constraints

SIF 3.0 Direct Environments are typically provided by an educational application which generates and maintains data that other applications need to access and / or update. Examples include systems as varied as an SIS, LMS, Portal or Data Store, all of which may want to provide clients with direct access to their supported data object types in a secure standardized way, without requiring the services of a middleware “Broker”.

In the effort to extend these applications to support the Direct Environments Provider interface, there is often the desire on the part of its developers to:

- Focus primarily on providing access to the application data in conformance with the particular SIF Data Model specification being implemented
- Reduce the level of support for the broad set of functionality defined in the SIF 3.0 Base Architecture document
- Reduce the number of Infrastructure and Utility Service operations which must be made available to its Environment Consumers.

The desire to minimize the functionality that such a service application has to support in order to interoperate directly with minimal SIF Consumers is contrasted with the equally important desire to maximize the Environment functionality any given Consumer developer can rely on to always be present.

The end result of this “push-pull” between contrasting goals is the set of mandatory Environment requirements detailed below. They both encompass and extend the Environment functionality utilized by the Minimal Consumer described in Appendix A.

This “mandatory” Environment interface addresses several important use cases:

- It allows any Service application such as an SIS or LMS which implements it to support a wide variety of directly connected SIF Client applications with little or no additional code required.
- It allows these SIF Client applications to register themselves as Consumers and run unchanged in any SIF 3.0 Environment, whether Minimal or otherwise.
- It minimizes the effort needed to write a SIF-compliant service application which can simultaneously implement BOTH:
 - One or more Data Object Service Providers interfaces in a Brokered Environment which are accessible to Consumers only through middleware
 - A Direct Environments Service Provider interface for its own set of Minimal Consumers, which can directly connect to it to obtain access to the same set of Data Object functionality.

B.2 Minimal Environment Architectural Requirements

The following list represents that portion of the Base Architecture functionality which **must** be supported by all Minimum Environments Providers.¹⁴

Functionality	Mandatory Support	Optional Support
Protocol Layer	HTTPS	Same + HTTP. Environment Consumers can attempt to utilize HTTP (via the initial URL used to create an Environment) and will succeed if HTTP is optionally supported by the Environment Provider.
Authentication Method	HMACSHA256	Same + BASIC. Environment Consumers can attempt to utilize BASIC Authorization and will succeed if BASIC Authorization is optionally supported by the Environment Provider. ¹⁵
# of Objects per Request	Both single and multiple	Same
Service Provider Type (ex: <i>Student</i>) Scope	Single Zone (default) contains all Services, and only a predefined set of required Data Model-specific Contexts for each Service Type (ex: " <i>Current</i> ")	Multiple Zones and Contexts per object type. Contexts can be dynamic, and be defined by Site Administration (ex: " <i>WilcoxJuniorHigh</i> ") as well as the Data Model.
Events and delayed Responses	Unsupported. There are no Events supported and all Responses are Immediate.	Support via inclusion of Subscriptions, Events Connector and Queues Infrastructure Services.
Queries	Paged Queries	Same + Service Paths, Dynamic Queries and predefined static XQuery Script functionality

¹⁴ This table will serve as a guideline for developing the SIF Infrastructure 3.0 Environment Provider Certification Suite. For further discussion of each of these areas, please refer to the Basic Architecture document.

¹⁵ While both BASIC and HMAC SHA256 are used to Authenticate Consumers (confirming its identity) common conventions refer to them as Authorization types, and the HTTP Header "Authorization" field will reflect that usage.

B.3 Core Infrastructure Service Support

The required set of core Infrastructure Service functionality which a Minimal Environments Provider **must** supply to its Consumers is described in the table below.¹⁶ It is contrasted with the (optional) higher level of functionality supplied by more featureful Environments Providers.

Note: Support is mandatory only for the Environments and Requests Connector Infrastructure Services.

Infrastructure Service(s)	Minimal Environment Infrastructure Services	Optional Environment Infrastructure Services
Environments, Environment and Provision Requests	<p>The authentication method enforced between Consumer and Environment is HMACSHA256, because the minimal Environment is likely to be provided by a web service.</p> <p>The <i>environments</i> Service in a Minimal Environment supports the <i>create</i> Environment operation, and should immediately return (on success) the Environment object data in the response.</p> <p>It also supports the <i>delete</i> Environment operation. There is no support for Environment <i>update</i> or <i>query</i>.</p> <p>All provisioning is pre-fixed and based upon the supplied Authentication information. There is no way for a Consumer to provision additional resources, and the <i>provisionRequests</i> Service is unavailable.</p> <p>If the appropriate provisioning is not contained in the Response to Environment creation (perhaps due to a lazy provisioning policy at the deployment site), the minimal Consumer should explicitly invoke needed operations and, if an Error Response indicates the operation was not authorized, <i>create</i> an Alert reporting that fact.</p>	<p>The BASIC authentication method may also be supported.</p> <p>Identical.</p> <p>The <i>environment</i> Service may also supports the “<i>query</i>” operation</p> <p>The <i>provisionRequests</i> Service essentially supports a controlled <i>update</i> to the contents of the Consumer’s Environment.</p> <p>Support of both the Zones and Providers Registries allows a more featureful Consumer to self-configure by discovering and provisioning itself to issue additional Service Provider operations.</p>
Requests Connector	<p>All requests for Service operations are invoked through this end point.</p> <p>There are no Connector-specific operations for the Consumer to request.</p>	Identical
Queues	Support for the Queues Infrastructure Service is	Queues Service support includes the ability to “ <i>create</i> ” a Queue Instance, which allows a

¹⁶ For further discussion of each of these Infrastructure Services, please refer to the Infrastructure Services document

	not required. The ability for a Consumer to access asynchronously arriving messages (Events or Delayed Responses) is not supported.	Consumer to request the oldest pending asynchronously arrived message.
Subscriptions	Support for the Subscriptions Infrastructure Service is not required. The ability for a Consumer to subscribe to and receive Service Events is not supported.	The Subscriptions Service <i>create</i> associates an allocated Queue Instance with the Events of a particular Service Provider, to capture all published Events by that Provider

B.4 Utility Service Interactions

The required set of Utility Service functionality which a Minimal Environments Provider **must** supply to its Consumers is described in the table below.¹⁷ It is contrasted with the (optional) higher level of functionality supplied by more featureful Environments Providers

Note: Support is mandatory only for the Alerts Utility Service.

Utility Service(s)	Minimal Environment Utility Services	Optional Environment Utility Services
Alerts	Any Consumer must be able to issue a <i>create</i> request to this Utility Service to create an Alert describing an unexpected error condition or interoperability mismatch it encountered on a Response to a previous Request. Only the <i>create</i> Alerts operation needs to be supported.	A Consumer can be given the right to <i>query</i> its own Alert objects
Zones Registry	Support for the Zones Registry is not required. There is only one Zone (the default) contained in a Simple Environment. The name of that Zone is the value of the “defaultZone” element returned after successful Registration. The Consumer in a Minimal Environment does not need to specify a Zone in any Request. The default one will be assumed.	The <i>zonesRegistry</i> Service contents can be fixed, and they are returned whenever a <i>query</i> operation is invoked. No other Service operations are standardly supported. At a minimum, only two Zones are needed. The first is the <i>environment-global</i> Zone which contains all available Utility Services, and the second is the Zone which contains all the Object Services available to all Consumers.
Providers	Support for the Providers Registry is not	Where it exists, there is one Providers Registry entry for every accessible Utility and Data Object

¹⁷ For further discussion of each of these Utility Services, please refer to earlier sections of this Utility Services document.

Registry	<p>required.</p> <p>The Consumer in a Minimal Environment must know the name of the type of every Object Service it will interoperate with, and if necessary, any associated Context other than <i>DEFAULT</i> (which is assumed where no Context is specified).</p>	<p>Service in the Consumer's Environment.</p> <p>For Direct Environments, the contents of this Registry is always fixed and returned whenever a <i>query</i> operation is invoked.</p> <p>For Brokered Environments, the <i>create</i> operation is available, and issued by any Consumer which is provisioning itself as a Service Provider.</p> <p>No other Service operations are standardly supported.</p>
Namespaces Registry	<p>Support for the Namespaces Registry is not required.</p> <p>The Consumer in a Minimal Environment must utilize the defined Infrastructure namespace, and stay within the set of additional Data Model namespaces understood by the Environment Provider.</p>	<p>Where it exists, all namespaces utilized by the Environments Provider (including company-specific ones) should have an entry in the Namespaces Registry or be included through a relative path in another XML Schema.</p> <p>The contents of this Registry can also be fixed and returned whenever a <i>query</i> operation is invoked.</p> <p>No other Service operations are standardly supported.</p>
Code Sets Registry	<p>Support for the External Codesets Registry is not required.</p> <p>The Consumer in a Minimal Environment must adopt the "blind trust" strategy, and assume any Code Set value it does not understand is legal, and accept it.</p> <p>The Environment Provider may also adopt this strategy, or it may reject any Request with a Code Set value it does not understand.</p>	<p>The Consumer (or Provider) can optionally attempt to verify whether an "unexpected" code set value it has received is simply unknown to it or illegal, by issuing a <i>query</i> to the Code Set Registry Service.</p> <p>No other Service operations are standardly supported.</p>
XQuery Templates Registry	<p>Support for the XQuery Templates Registry is not required</p> <p>The Consumer in a Minimal Environment can only qualify a <i>query</i> request by inserting a Basic XQuery script as one of the Query Parameters in the URL of the Service to which it is issuing the Query</p>	<p>If the Service Providers in an Environment are capable of processing a set of "static" XQuery templates (with the Template Parameters provided as Query Parameters in the Service URL which the Consumer used as the destination of the Query), the <i>query</i> operation on this Registry should be supported.</p> <p>If XQuery Templates can be dynamically created by Consumers, both <i>create</i> and <i>delete</i> operations should be supported.</p>

		Due to template caching, the <i>update</i> Service operation is not standardly supported.
--	--	---

Appendix C: Minimal Service Providers

Service Providers operate in Brokered Environments only, and are limited to those Service Consumers who are additionally authorized to declare themselves to be Providers for one or more objects.

The Consumer's initial authorization rights **must** include the ability to declare itself as a Provider of one or more object types ... or it should be able to obtain such access through a successful *create* issued to the *Provision Requests* Infrastructure Service by which means it **may** explicitly assert such authorization.

Whether or not it actually issues that create, the prospective Service Provider can then declare itself to be a Provider of one or more given Application Services (each with a specified Context within a specified Zone) by successfully issuing one or more *create* Requests (single or multiple) to the Providers Registry Utility Service. A simple example of such a multiple request would be when a Consumer is asserting the right to provide Student objects for Consumer applications spanning more than one Zone in the Environment. Each Provider Registry Entry **must** include a unique URL to which authorized requests from other Consumers for that specific Service will be dispatched.

If successful, a Providers Registry Event will be issued announcing the availability of the new service or services, and subsequently, Consumer Requests made for that Service will begin to arrive on the supplied URL, containing the Session Token of the Provider rather than the issuing Consumer (preserving Consumer security and allowing the Provider to determine that the Request was from a trusted intermediary – in this case the Message Broker). The Service Provider must respond to these requests immediately (i.e. on the HTTP/S Response that corresponds to the HTTP/S Request), whether or not:

- The Request involved one or more than one object
- The Response was successful or Erroneous
- The Consumer requested an immediate or delayed response

When change Requests (or internal updates) alter the content of one or more objects the Service is providing, it **must** issue a corresponding Event to the Event Connector Infrastructure Service (the URL of which **must** have been returned in its Environment when the Consumer initially registered).¹⁸

The Event Connector will determine the exact set of approved subscribers, and (assuming that any exist), route the Event to the indicated persistent FIFO Queue Instance specified by the each Requestor in its Subscription, transparently to the Service Provider.

Essentially the Minimum Service Provider is required to do only the following above and beyond what a Service Consumer would do:

- Declare itself a Provider by issuing an appropriate Create Request to the Providers Registry Utility Service containing a URL for incoming Service Requests, and hope for the best.
- If successful, accept and process arriving Requests for its Service(s) and respond “immediately”.

¹⁸ Note that this URL might not be made available to any Consumer which has not been pre-approved to be a Service Provider

- When data changes occur, publish the corresponding Event to the Event Connector.